

Operating System Concepts

A. Silberschatz, P. B. Galvin, and G. Gagne, 10th edition

Chapter 4: Threads & Concurrency

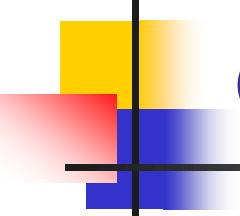
Mei-Ling Chiang, Professor

**Dept. of Information Management
National Chi-Nan University, Taiwan, ROC**



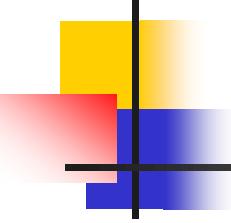
- 
- A **process** is defined by the **resources** it uses and by the **location** at which it is executing.
 - Processes serve 4 roles:
 - Unit of resource allocation
 - Memory context
 - Program that is being executed
 - Unit of activity
 - What do you do if your program wants to do more than one thing at a time ?

- 
- The process model introduced in Chapter 3 assumed that a process was an executing program with a single thread of control.
 - Modern OSes now provide features for a process to contain multiple threads of control.
 - If a process has multiple threads of control, it can perform more than one task at a time.



Chapter Objectives

- Identify the basic components of a **thread**, and contrast threads and processes
- Describe the benefits and challenges of designing multithreaded applications
- Illustrate different approaches to implicit threading including thread pools, fork-join, and Grand Central Dispatch
- Describe how the Windows and Linux OSes represent threads
- Design multithreaded applications using the Pthreads, Java, and Windows threading APIs



Chapter 4: Threads & Concurrency

Outline

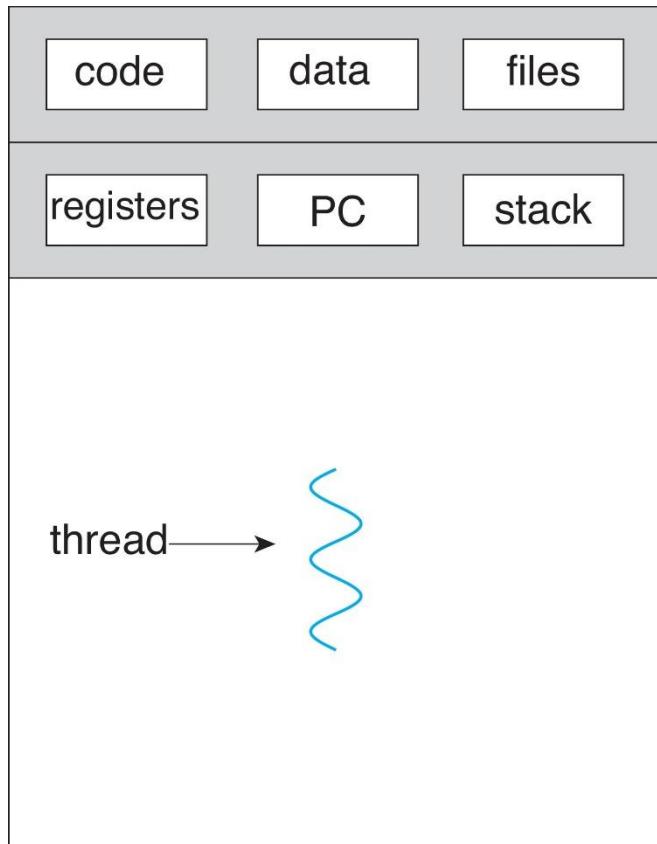
- **4.1 Overview** 
- **4.2 Multicore Programming** 
- **4.3 Multithreading Models** 
- **4.4 Thread Libraries** 
- **4.5 Implicit Threading** 
- **4.6 Threading Issues** 
- **4.7 Operating-System Examples** 



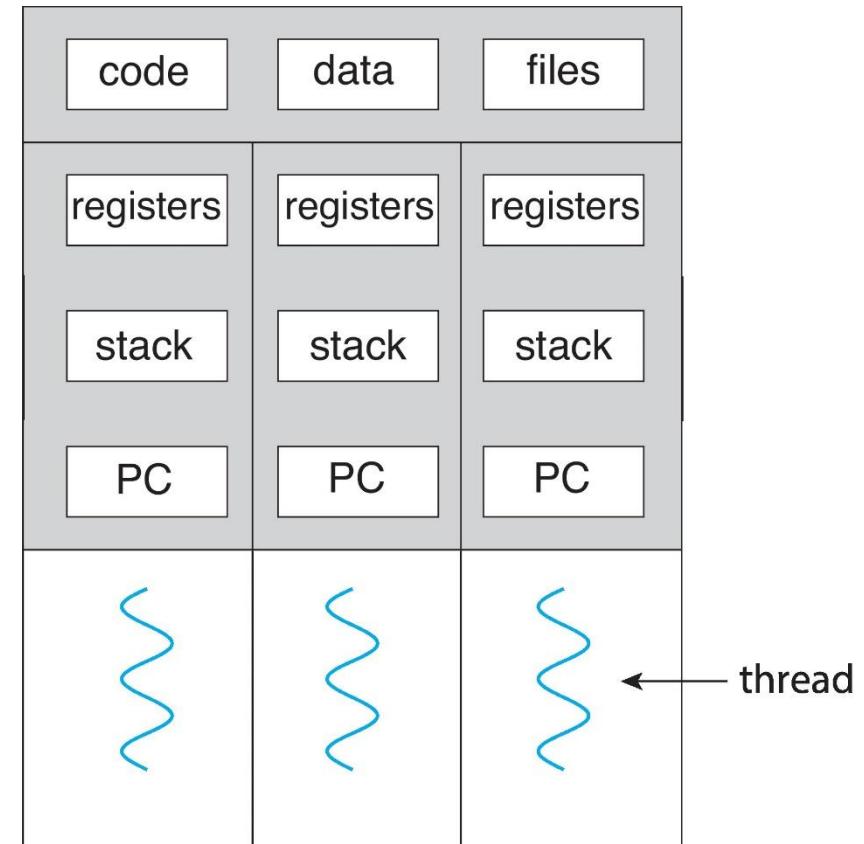
4.1 Overview

- A **thread** is a basic unit of CPU utilization, comprises
 - a thread ID, a program counter, a register set, and a stack
- A thread shares with other threads belonging to the same process its
 - code section
 - data section
 - other OS resources (open files, signals)
- A **traditional process** has a single thread of control. If a process has multiple threads of control, it can perform more than one task at a time.

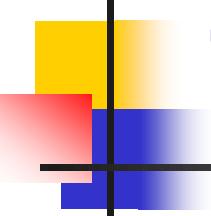
Single-threaded and multithreaded processes (Fig. 4.1)



single-threaded process

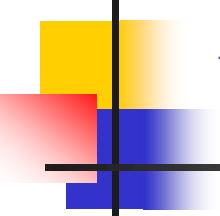


multithreaded process



4.1.1 Motivation

- **Most software applications that run on modern computers and mobile devices are multithreaded.**
- **An application typically is implemented as a separate process with several threads of control.**
- **Example: a web browser**
 - a thread displays images or text while another thread retrieves data from the network
- **Example: a word processor**
 - a thread for displaying graphics, a thread for responding keystrokes from the user, and a thread for performing spelling and grammar checking in the background
- **Example:**
 - applications can be designed and perform several CPU-intensive tasks in parallel on multicore systems.

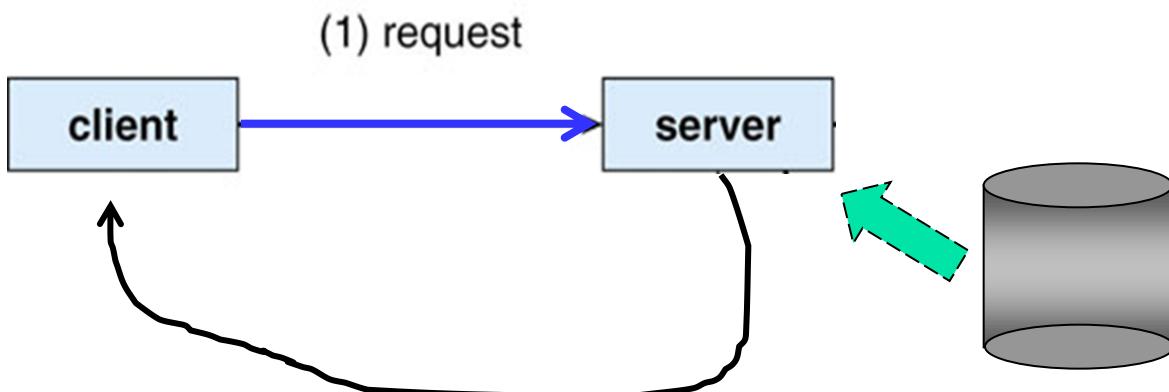


Motivation (Cont.)

- **In certain situations a single application may be required to perform several similar tasks.**
- **Example: web server**
 - Single-threaded process: service only one client at a time
 - Multiple processes
 - multi-threaded process
 - One process that contains multiple threads to serve the same purpose.
- **Process creation is time consuming and resource intensive.**
- **If the new process will perform the same tasks as the existing process, it is generally more efficient for one process that contains multiple threads to serve the same purpose.**

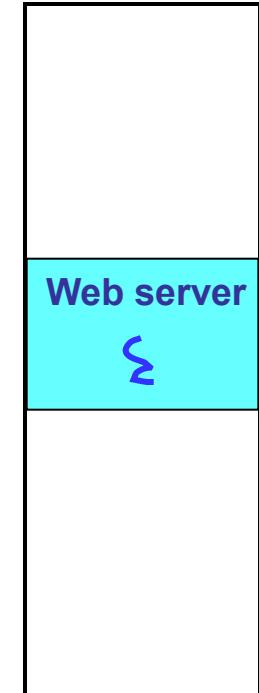
Example: Web Server

- Single-threaded process
 - service only one client at a time



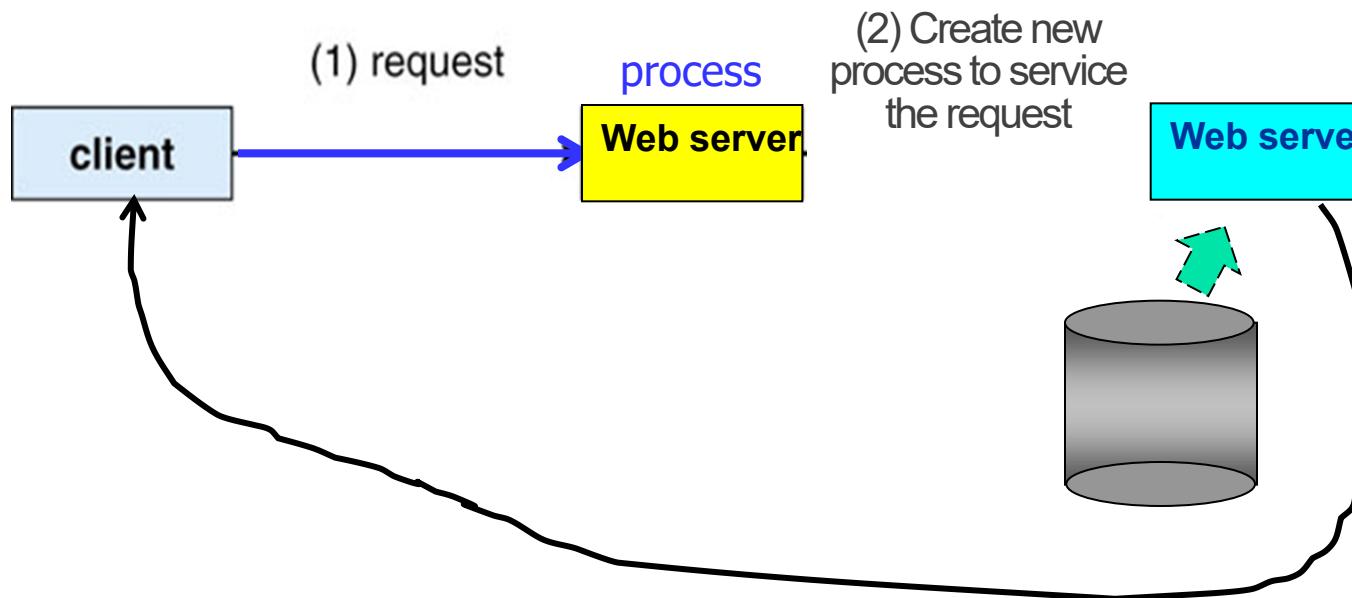
```
void main() {  
    while (1)  
    {  
        receive();  
        disk read;  
        send();  
    }  
}
```

RAM



Example: Web Server

■ Multiple processes



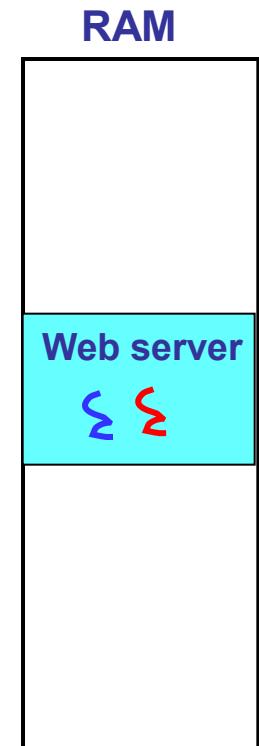
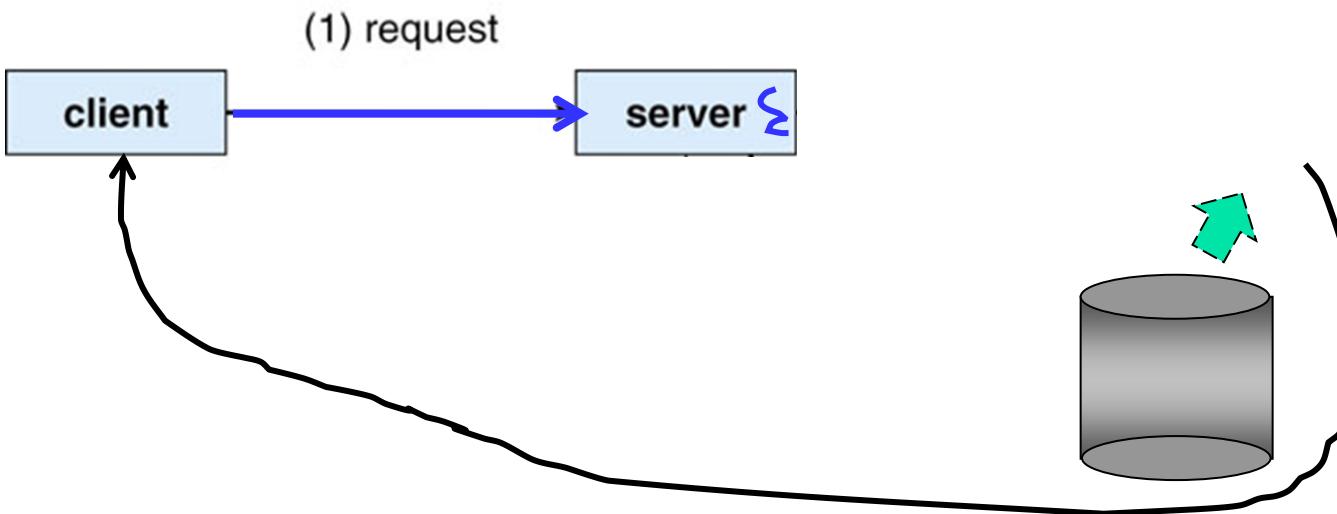
```
void main() {  
    while (1)  
    {  
        receive();  
        pid=fork();  
        if (pid==0) {  
            disk read;  
            send();  
            exit(0);  
        }  
    }  
}
```

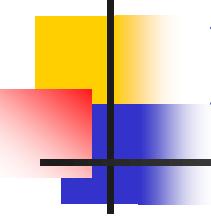


Multithreaded server architecture

- ***Multithreading a web server process***

```
void fun() {  
    disk read;  
    send();  
    thread_exit(0);  
}  
void main() {  
    while (1)  
    {  
        receive();  
        createthread(fun);  
    }  
}
```



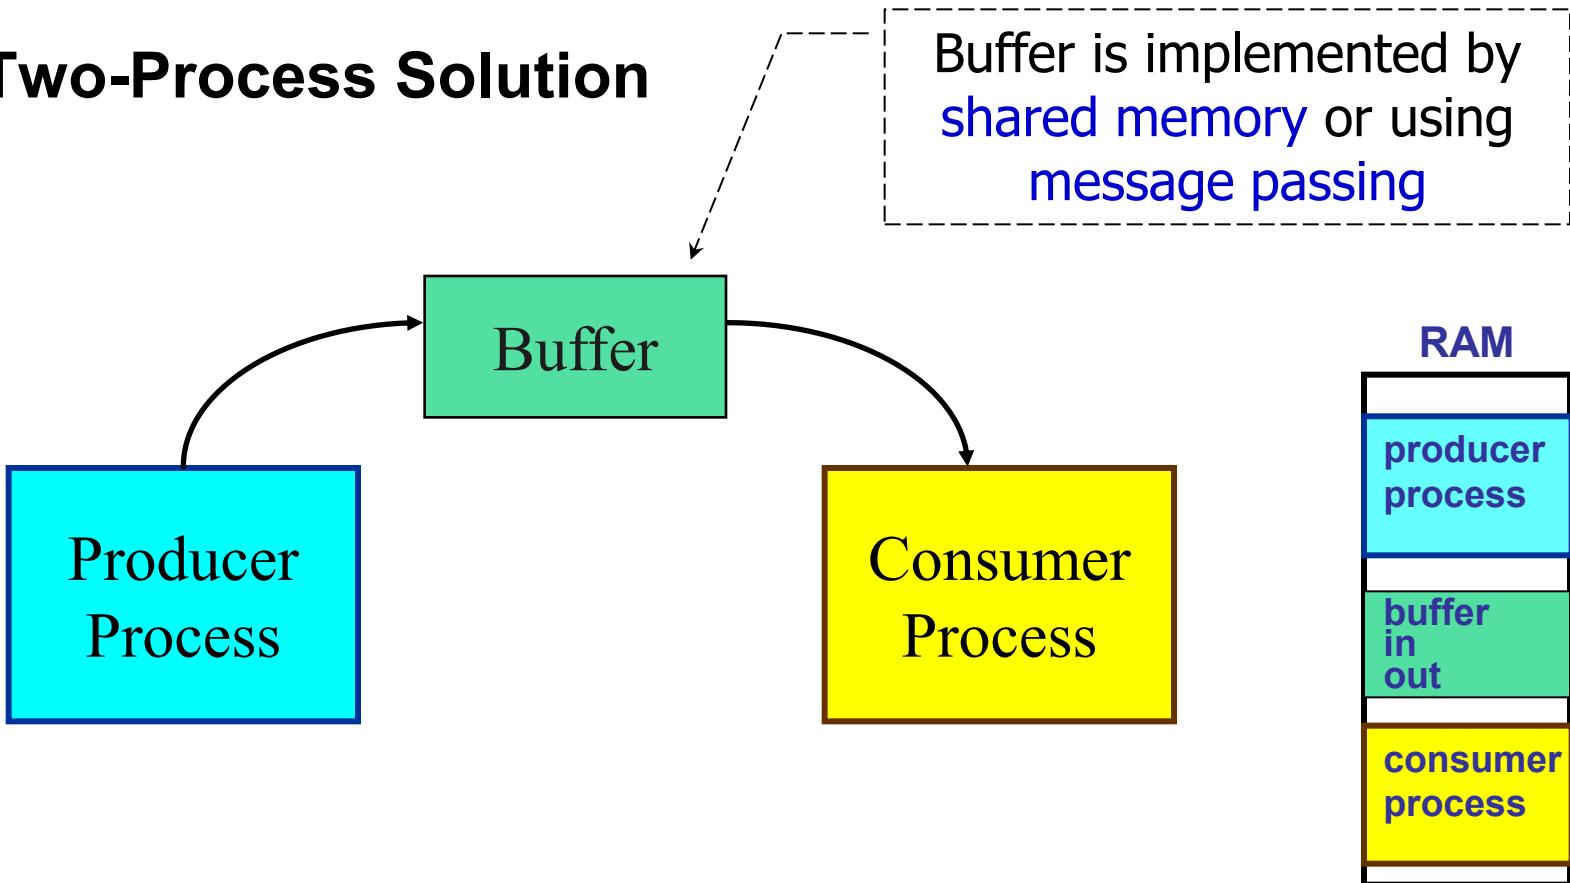


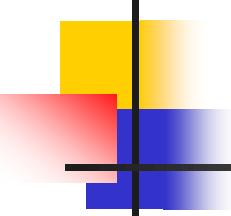
Motivation (Cont.)

- In a multiple threaded process, while one thread is blocked and waiting, a second thread in the same task can run.
 - Cooperation of multiple threads in the same job confers higher throughput and improved performance. 紹予
 - Applications that require sharing a common buffer (i.e., producer-consumer) benefit from thread utilization.

A Buffering Problem

- **Two-Process Solution**

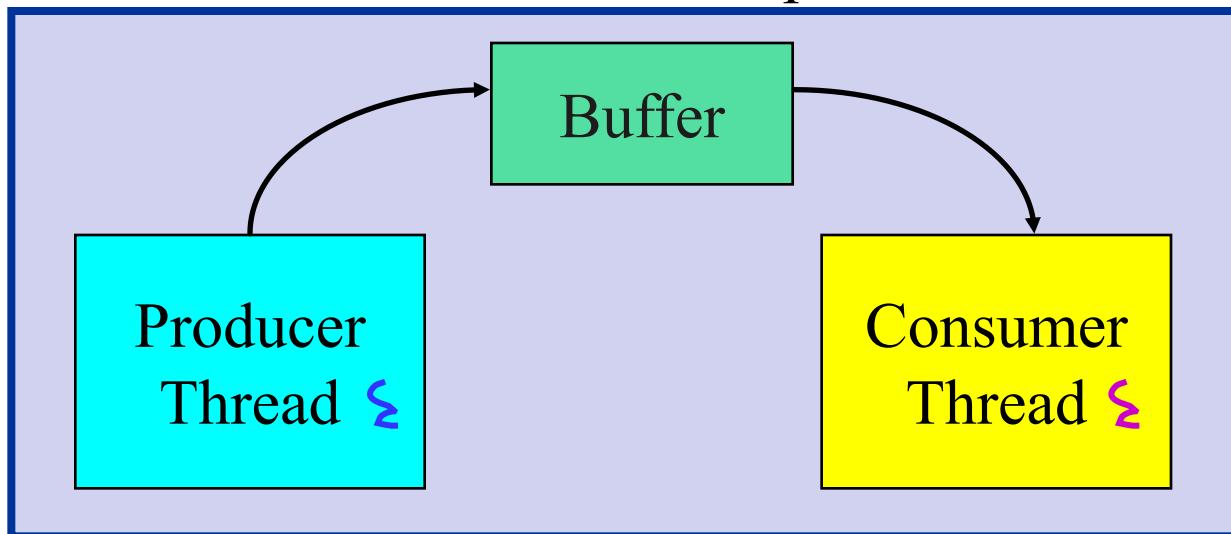


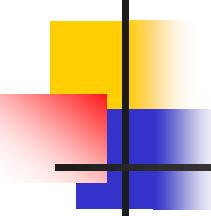


A Buffering Problem

- Two-Thread Solution

One process with 2 threads





Motivation (Cont.)

- **Most OS kernels are also typically multithreaded.**
- **Example:** During boot time on Linux systems, several threads are created. Each thread performs a specific task, such as managing devices, managing memory, or interrupt handling.

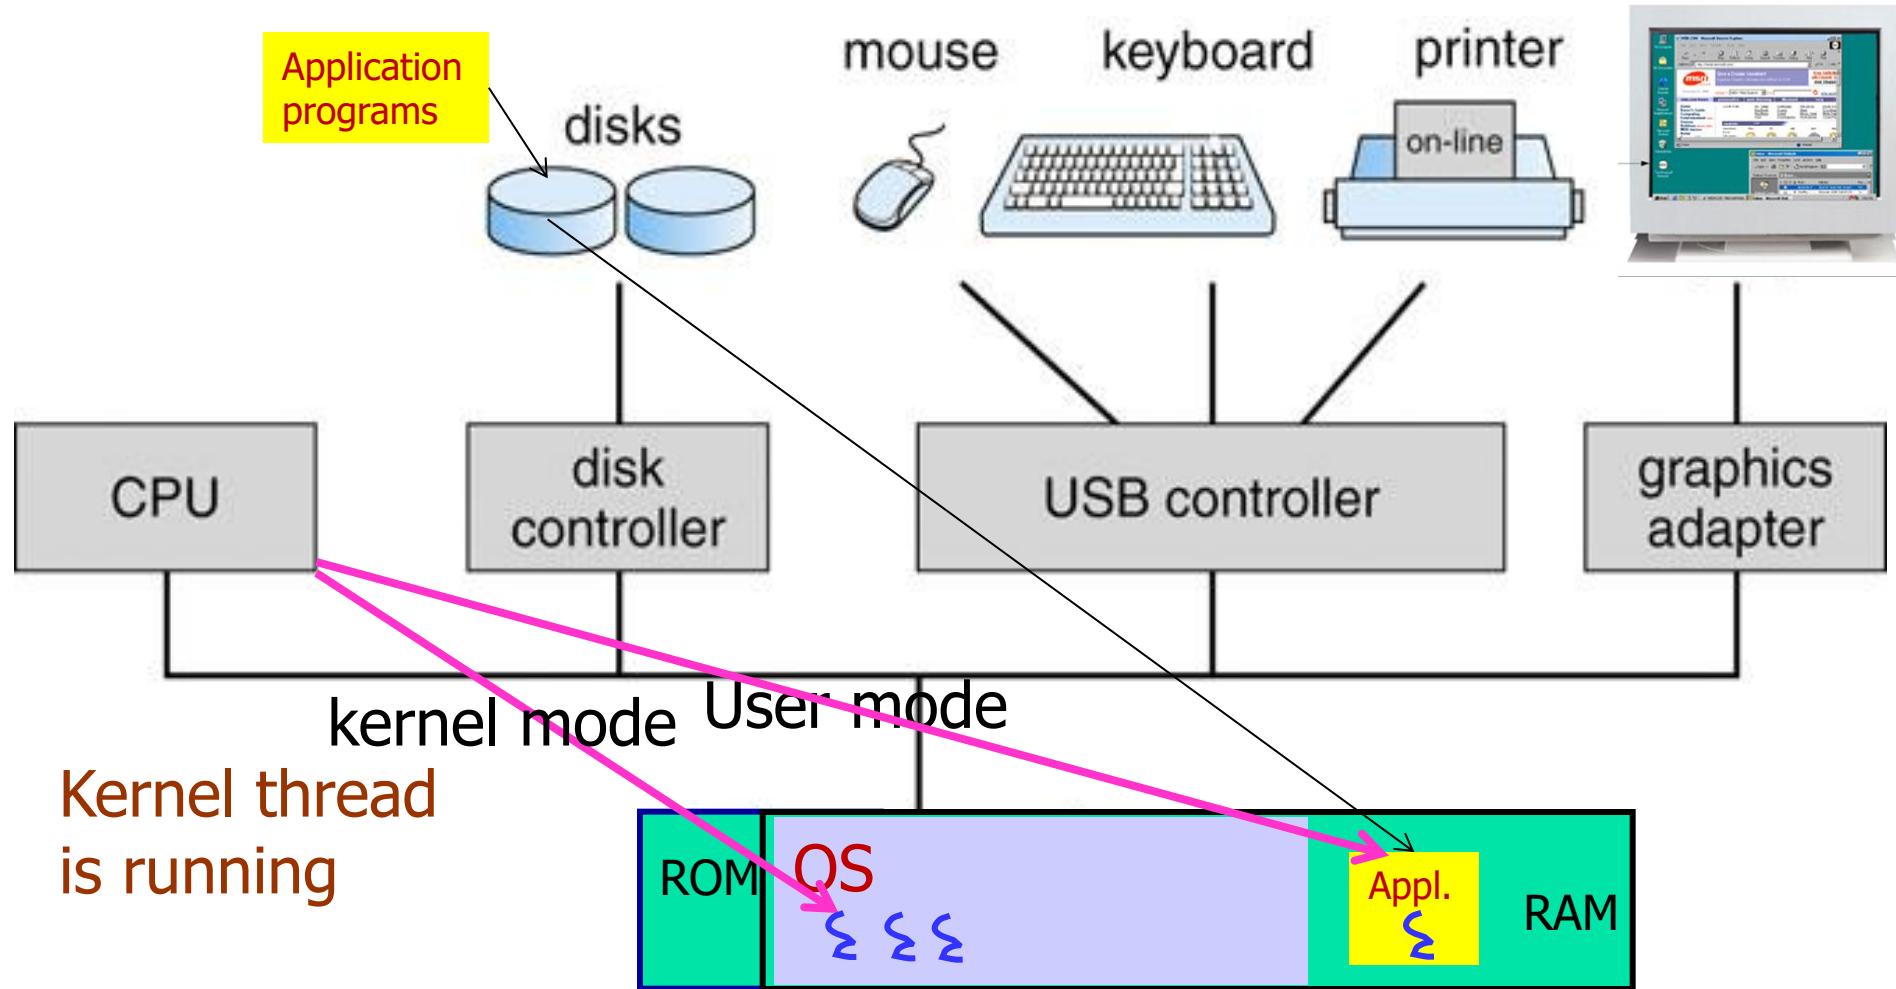
User applications

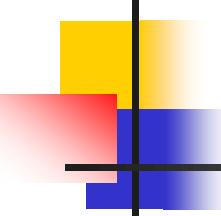
kernel
threads
ΣΣΣ

OS kernel

When does an OS run?

(5) kernel thread

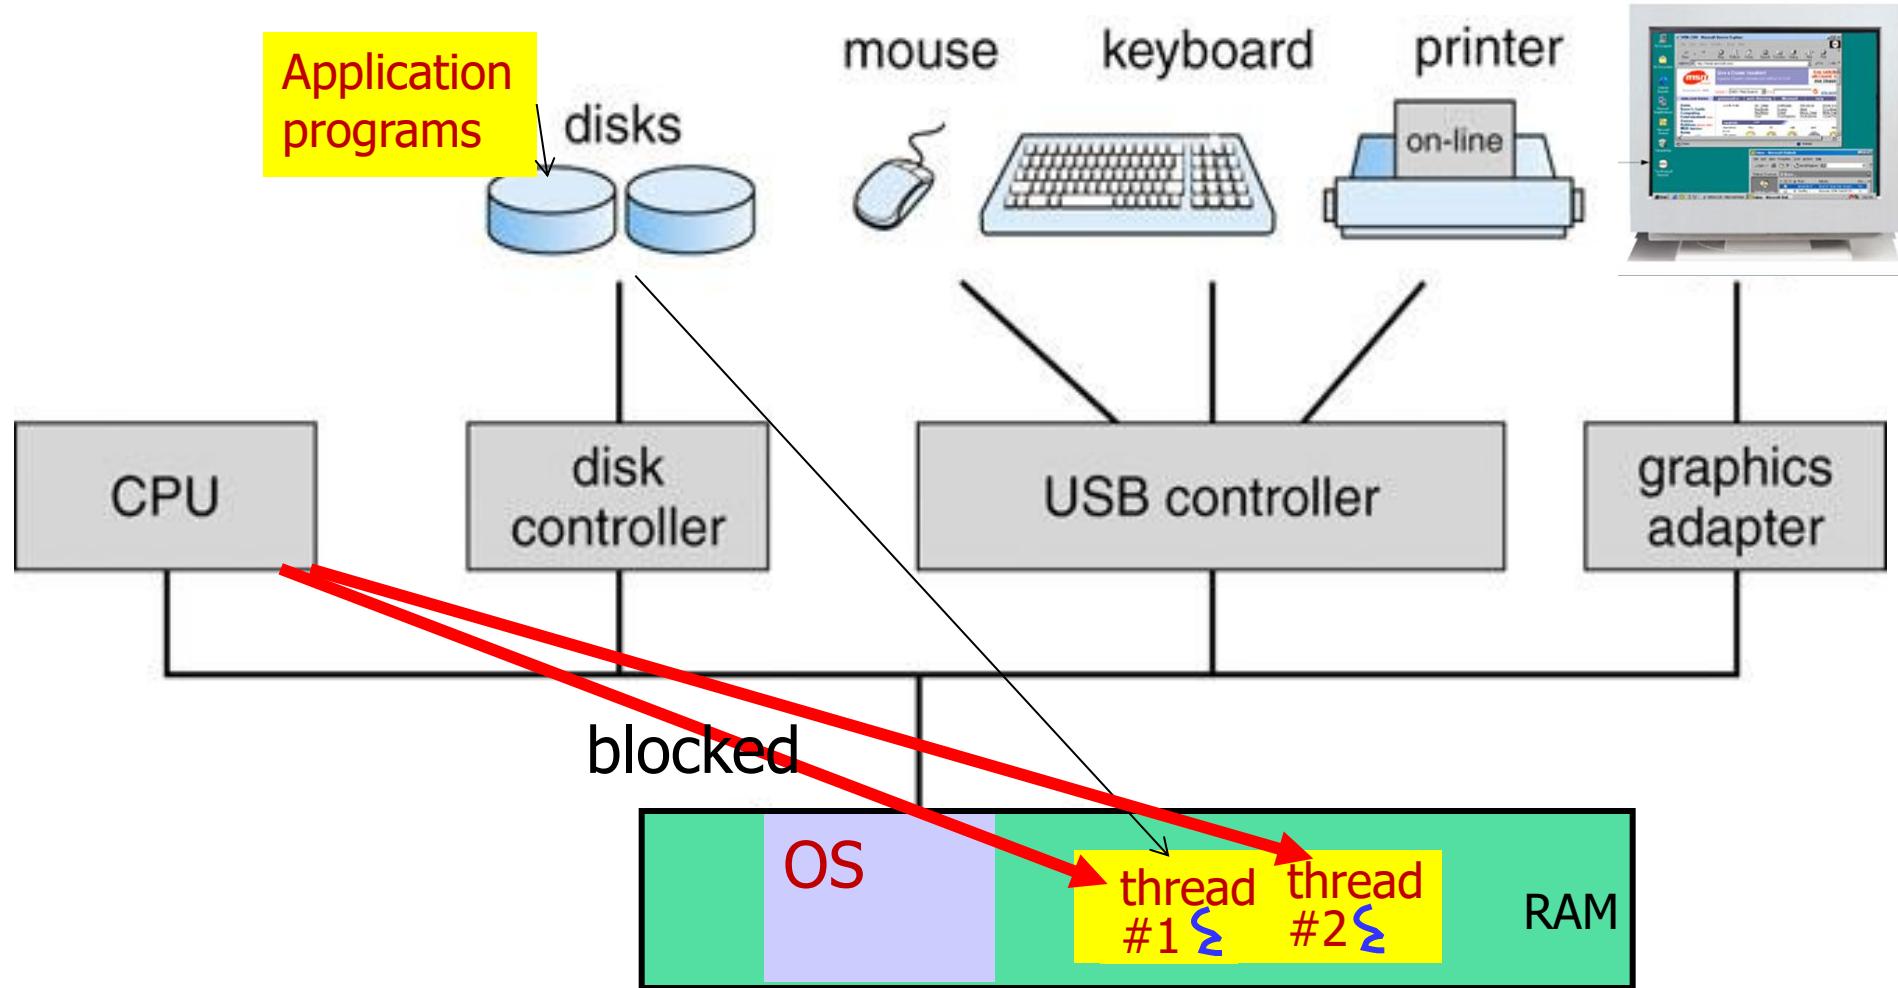




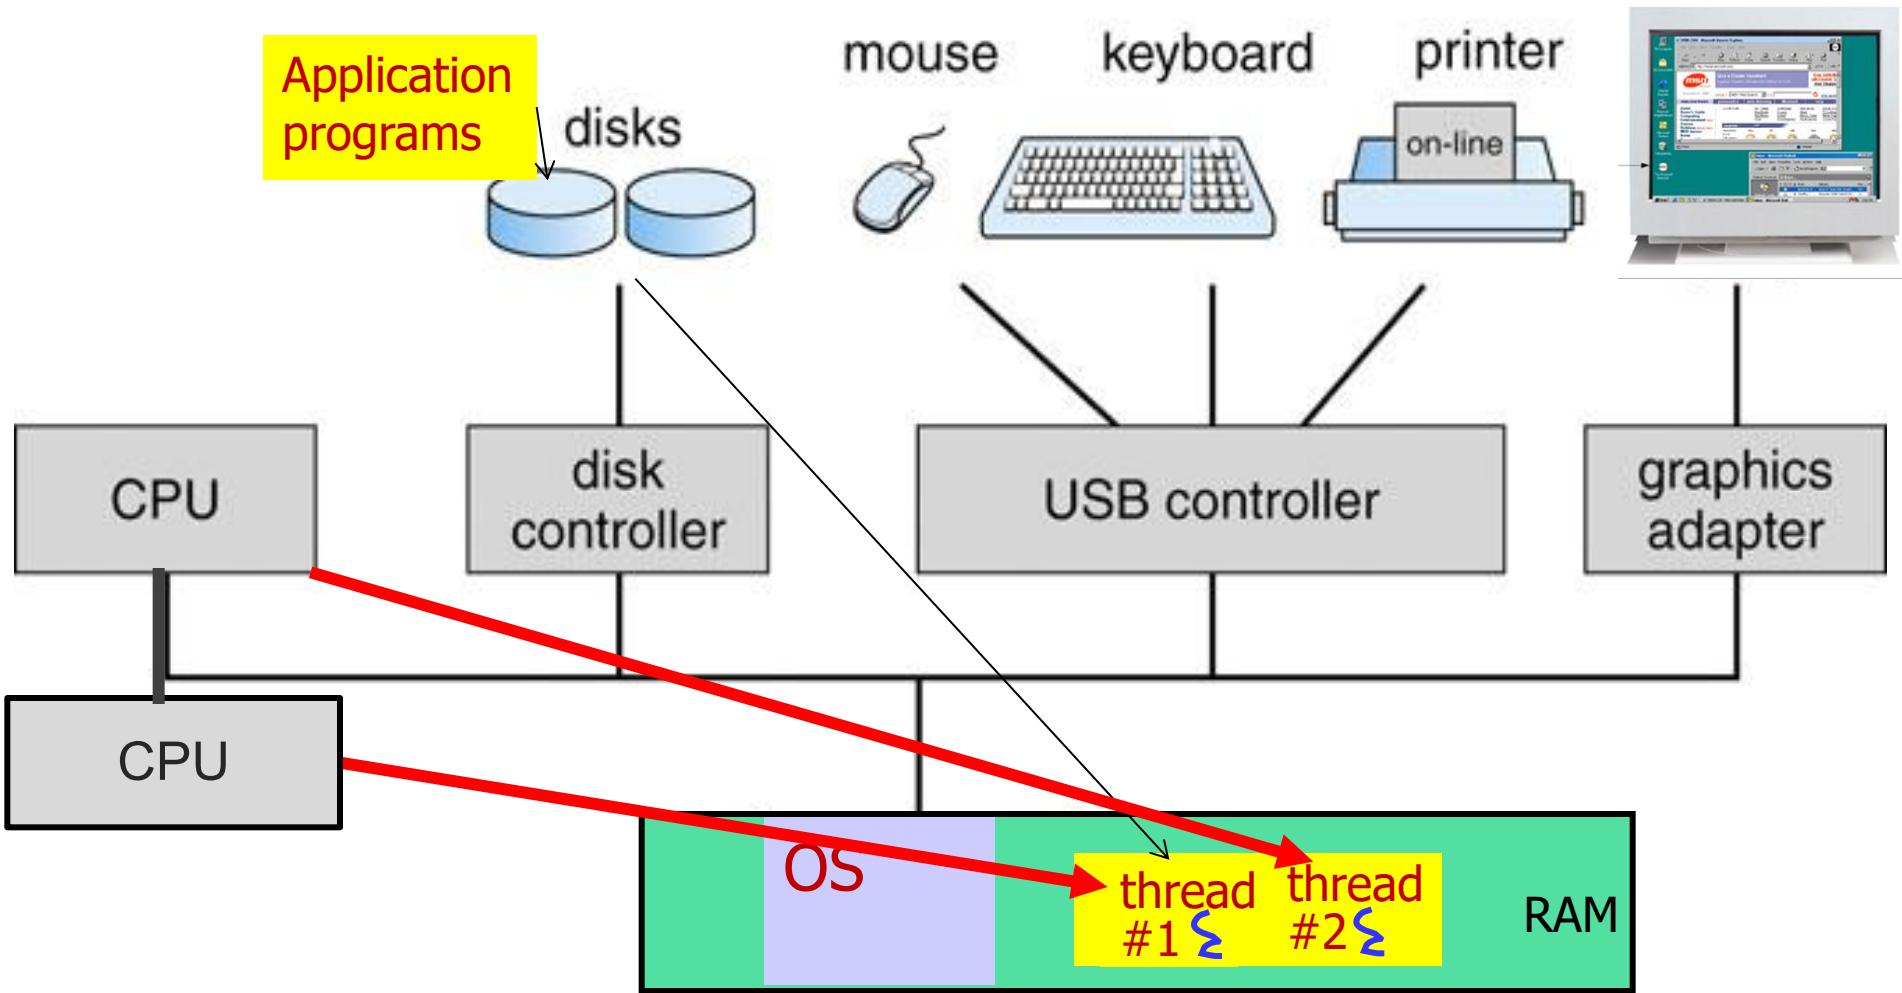
4.1.2 Benefits

- **The benefits of multithreaded programming:**
 - **Responsiveness**
 - Multithreading an interactive application may allow a program to continue running even if part of it is blocked or is performing a lengthy operation, thereby increasing responsiveness to the user.
 - **Resource Sharing**
 - Threads share the memory and the resources of the process to which they belong.
 - **Economy**
 - More economical to create and context-switch threads.
 - **Scalability**
 - can take advantage of multiprocessor architectures, where threads may be running in parallel.

Increase Responsiveness



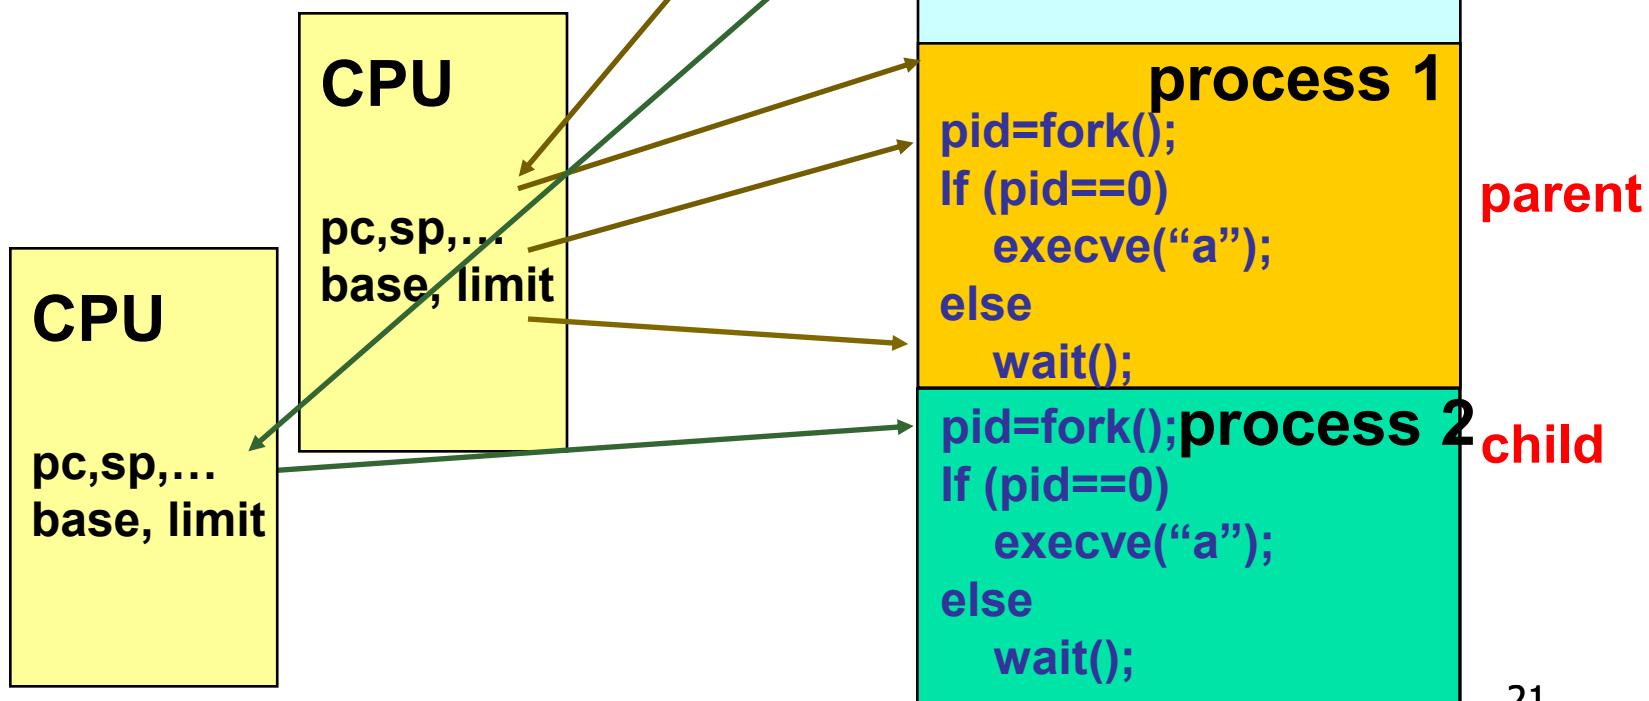
Parallel Execution



fork()

create a process in UNIX

CPU Scheduling



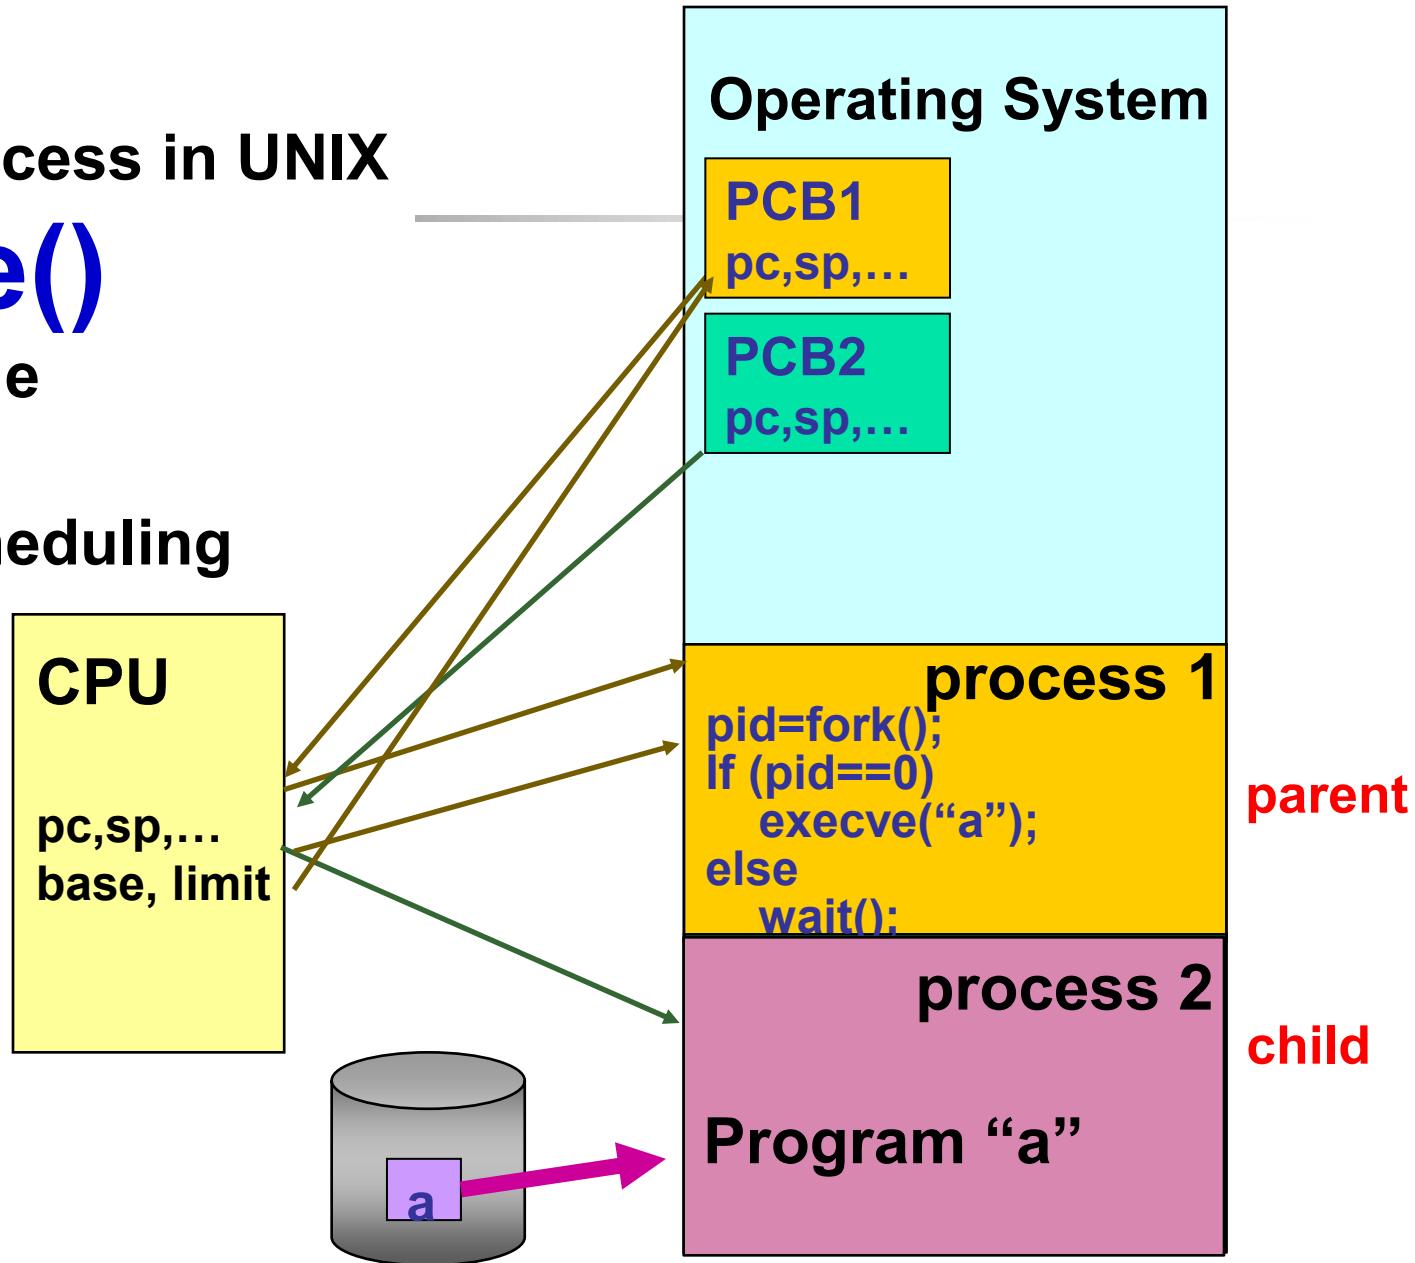
fork()

create a process in UNIX

execve()

execute a file

CPU Scheduling

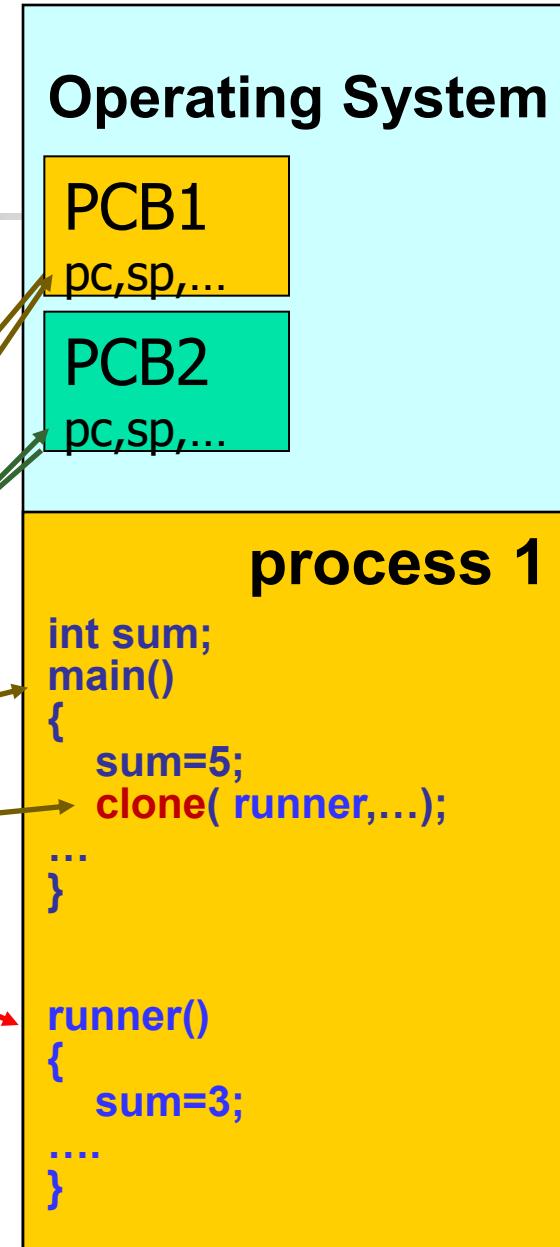
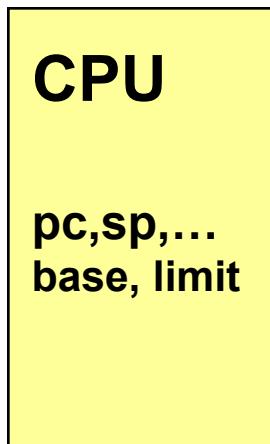
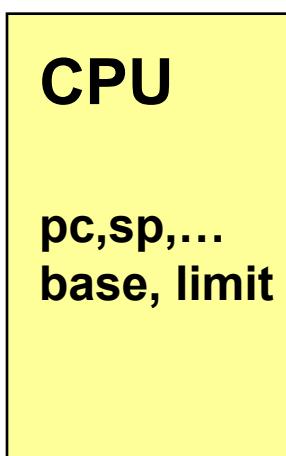


RAM

clone() a thread in Linux

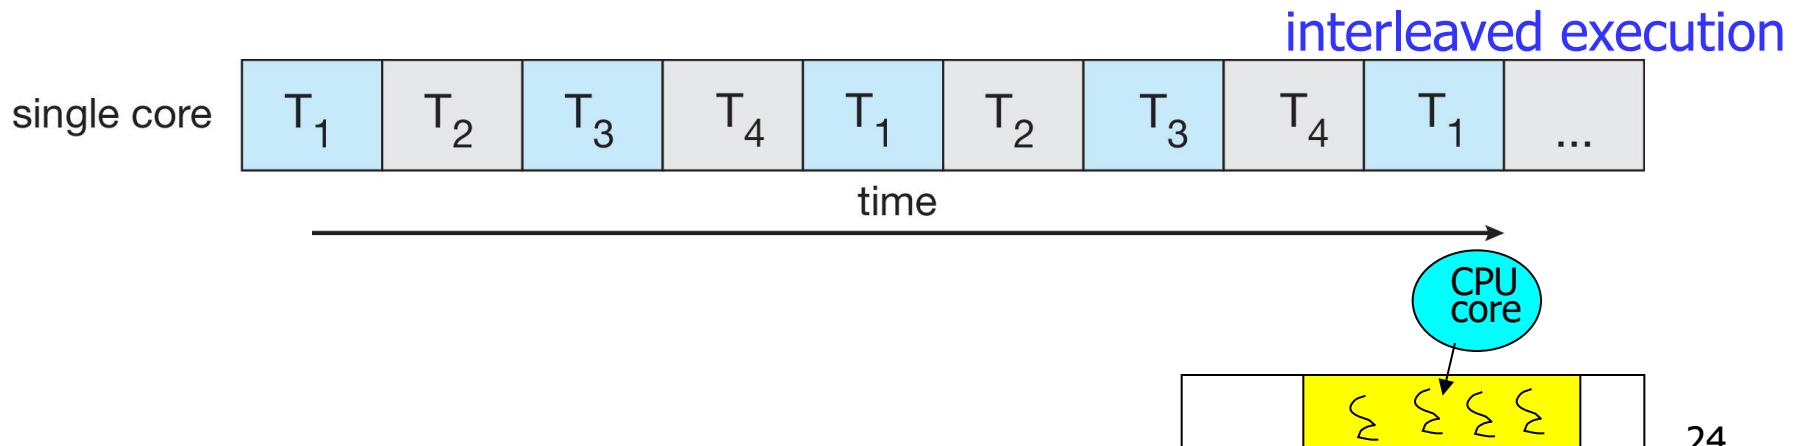
thread: a lightweight process

CPU Scheduling

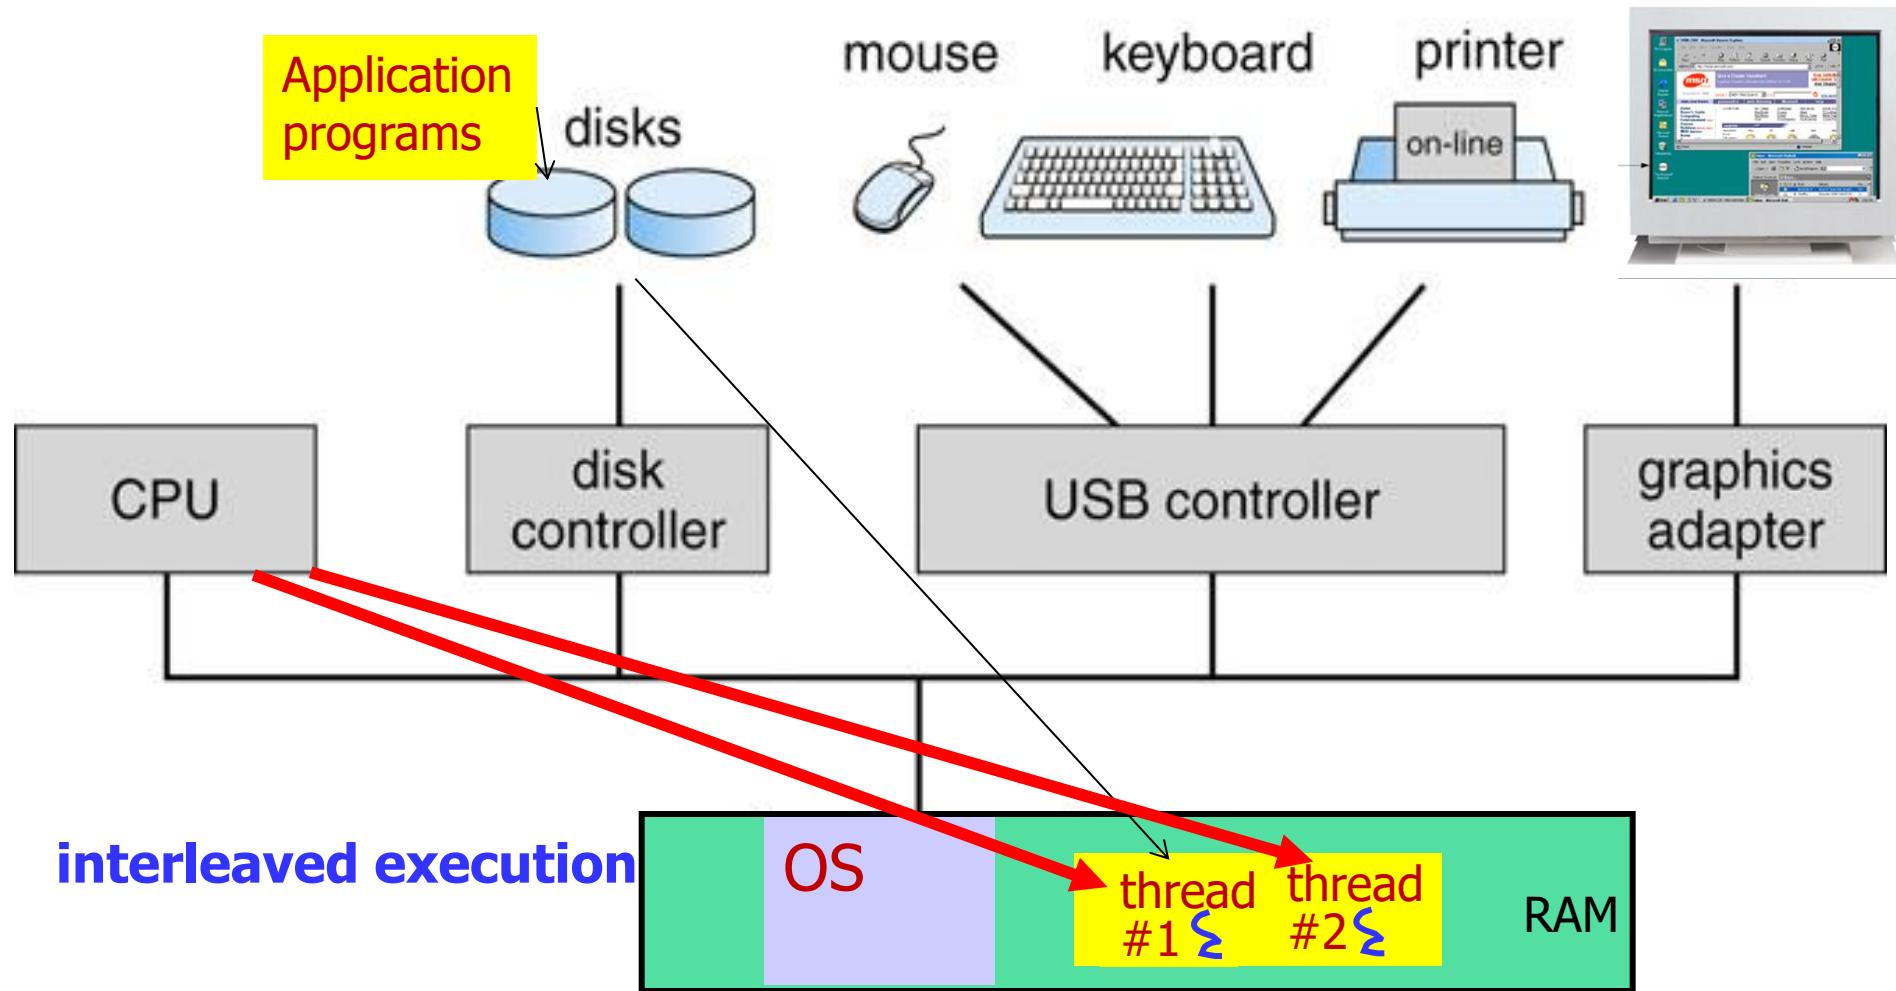


4.2 Multicore Programming

- Multithreaded programming provides a mechanism for more efficient use of multiple computing cores and improved concurrency.
- On a system with a single computing core, 同時發生 concurrency merely means that the execution of the threads will be interleaved over time.
交錯
- Concurrent execution on single-core system:



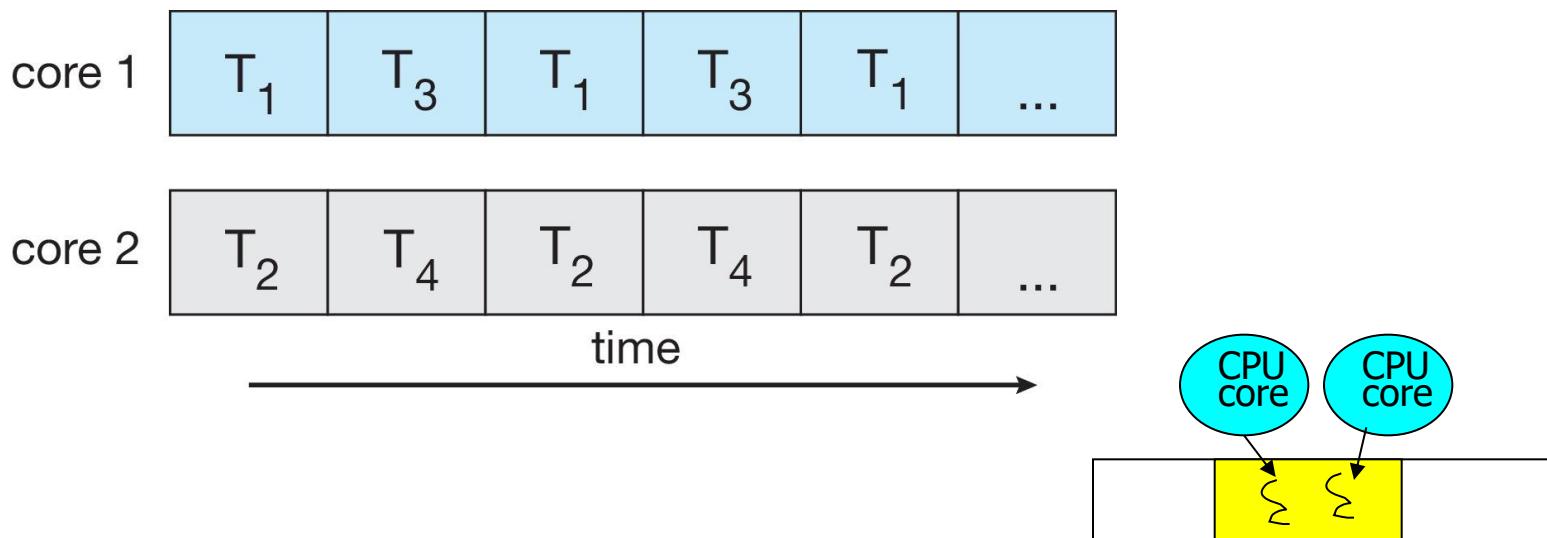
Concurrent Execution on a Single-Core Processor



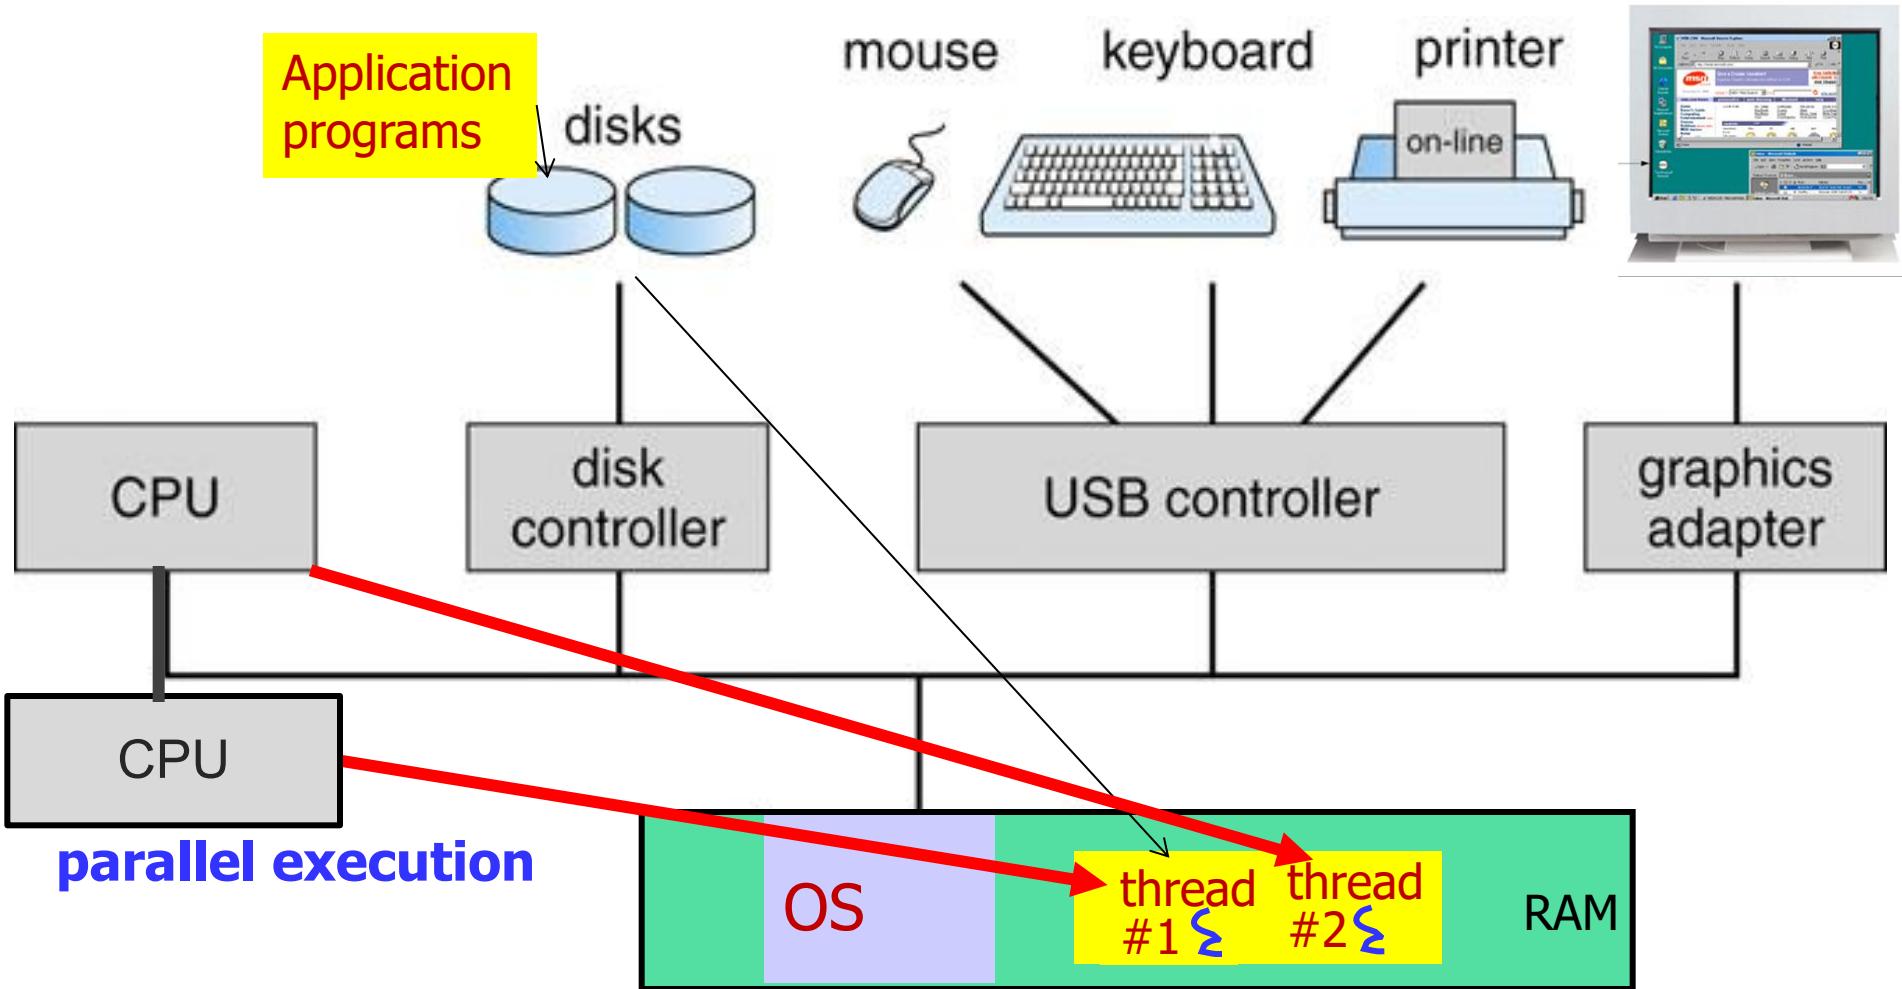
Multicore Programming (Cont.)

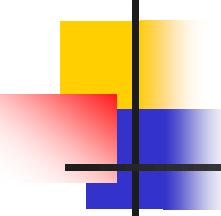
同時發生

- On a system with multiple cores, **concurrency** means that the threads can run in parallel, as the system can assign a separate thread to each core.
- **Parallel execution on a multi-core system:**



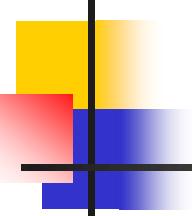
Parallel Execution





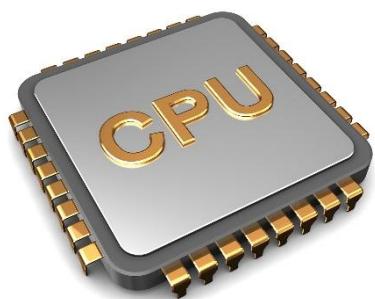
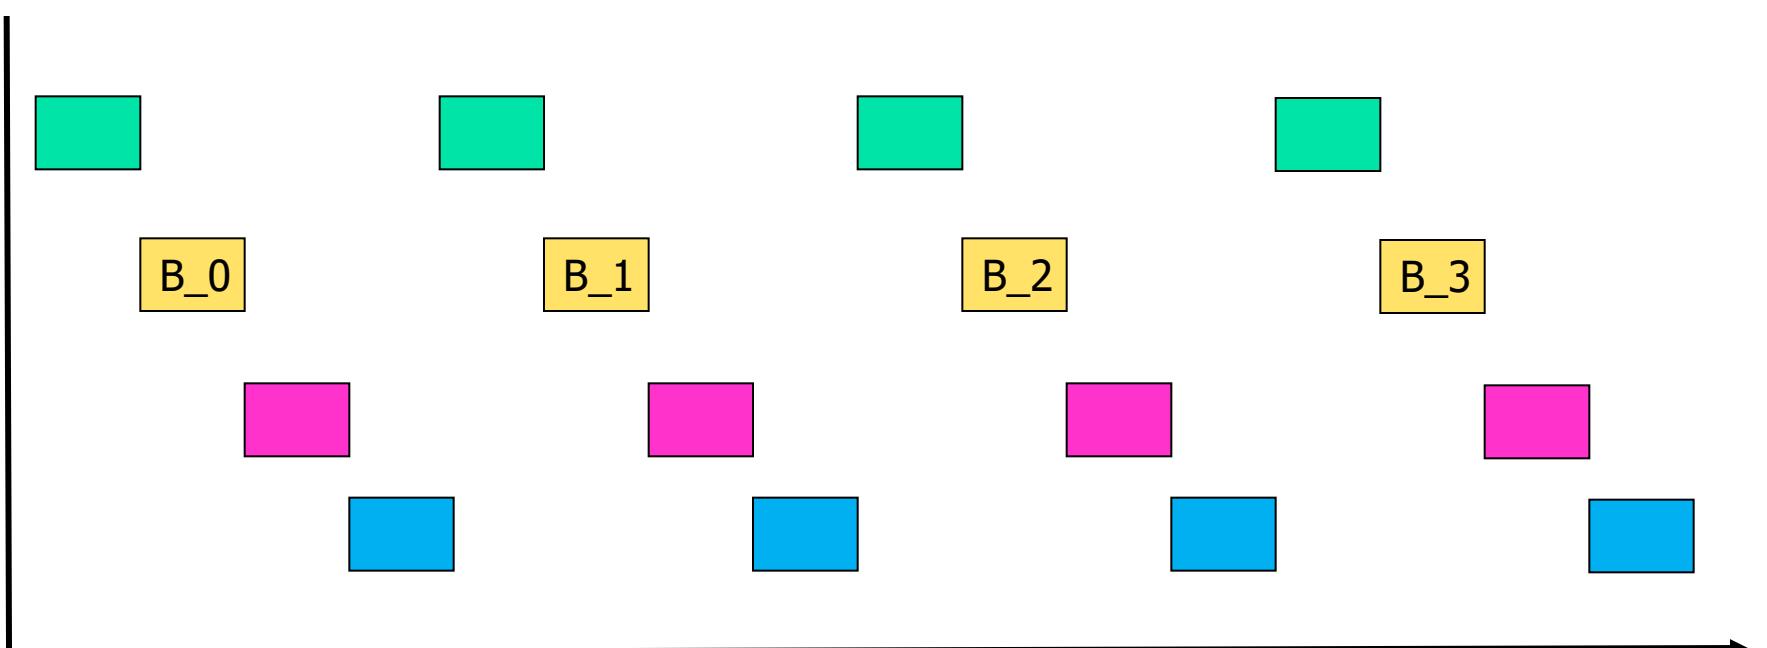
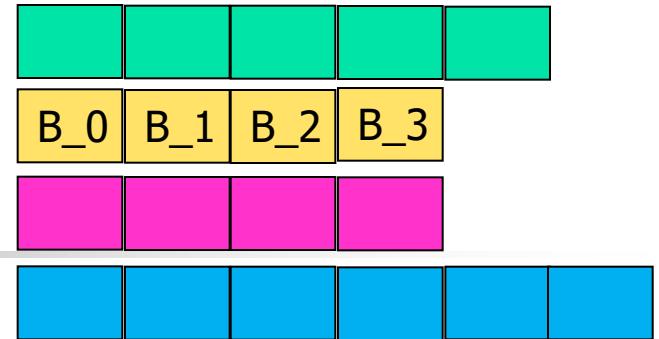
Concurrency vs. Parallelism

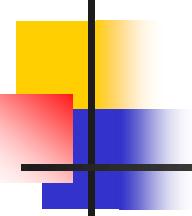
- A system is ***parallel*** if it can perform more than one task simultaneously
- A ***concurrent*** system supports more than one task by allowing all the tasks to make progress
 - Single processor / core, CPU scheduler providing concurrency
 - Provide the illusion of parallelism by rapidly switching between processes, thereby allowing each process to make progress
 - Processes were running concurrently, but not in parallel



Concurrency

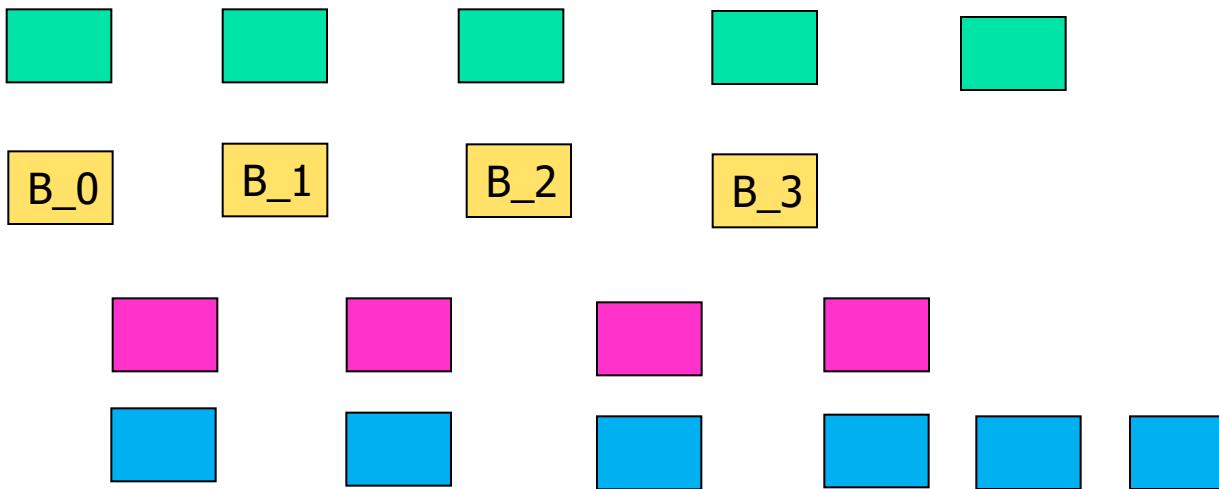
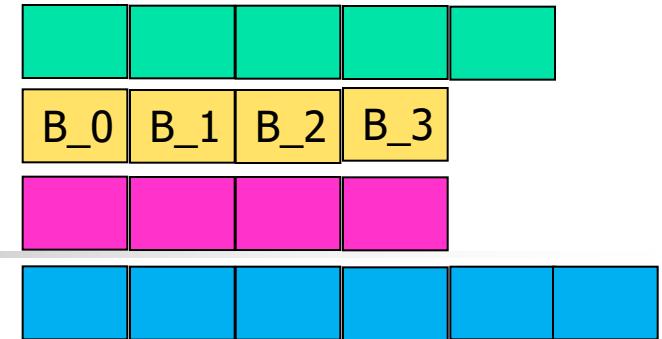
Tasks



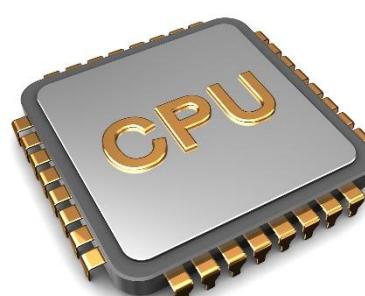
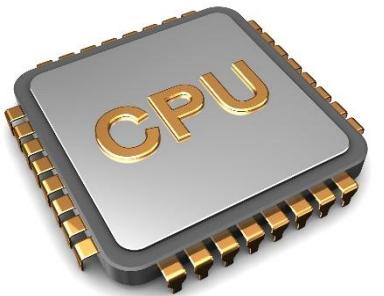


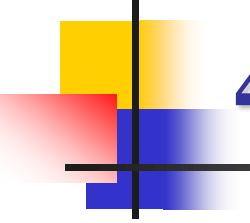
Concurrency

Tasks



Time

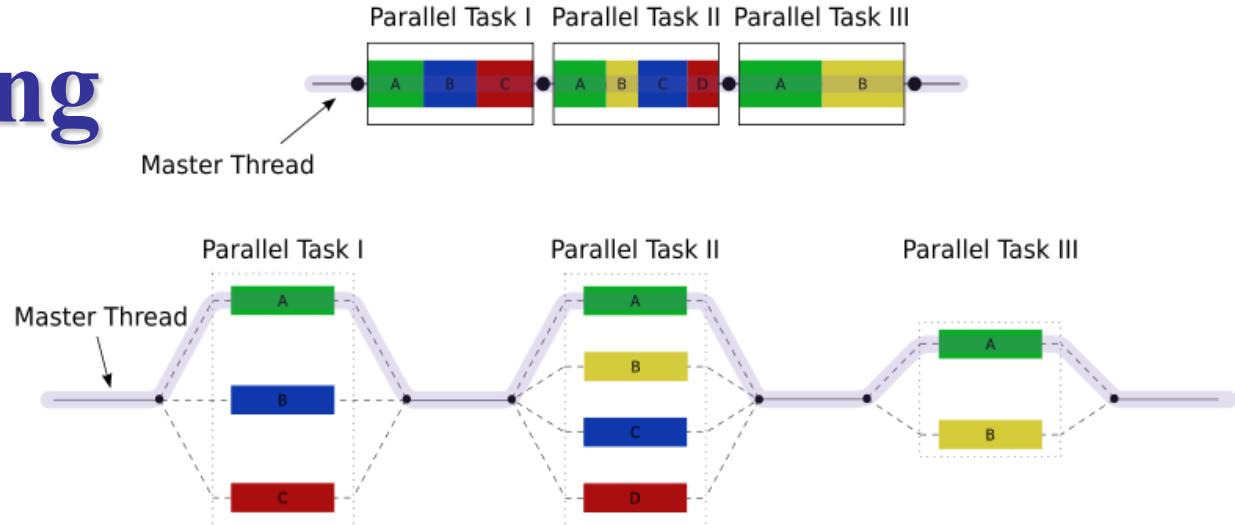




4.2.1 Programming Challenges

- **Multicore systems put pressure on system designers and application programmers to make better use of multiple computing cores.**
 - **OS designers:** write **scheduling algorithms** to allow the parallel execution on multiple processing cores
 - **Application programmers:** modify existing programs and design multithreaded programs

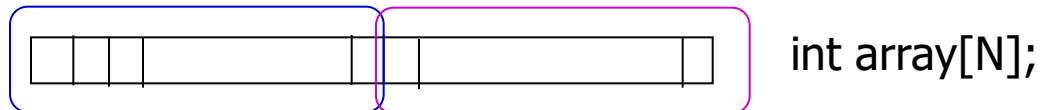
Programming Challenges (Cont.)



- **Challenges in programming for Multicore systems:**
 - **Dividing activities**
 - Identifying tasks that can run in parallel
 - **Balance**
 - Ensure that the tasks perform equal work
 - **Data splitting**
 - Data accessed by the tasks must be divided to run on separate cores
 - **Data dependency**
 - Ensure that the tasks' execution is synchronized to accommodate the data dependency
 - **Testing and debugging**

4.2.2 Types of Parallelism

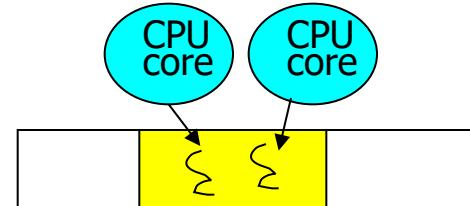
- **Data parallelism** – distributes subsets of the same data across multiple computing cores, perform the same operation on each core
 - e.g., two threads sum array[N] on a dual-core system in parallel

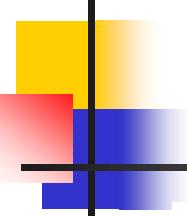


- thread A sums array[0]..[N/2-1] on core 0
- thread B sums array[N/2]..[N-1] on core 1

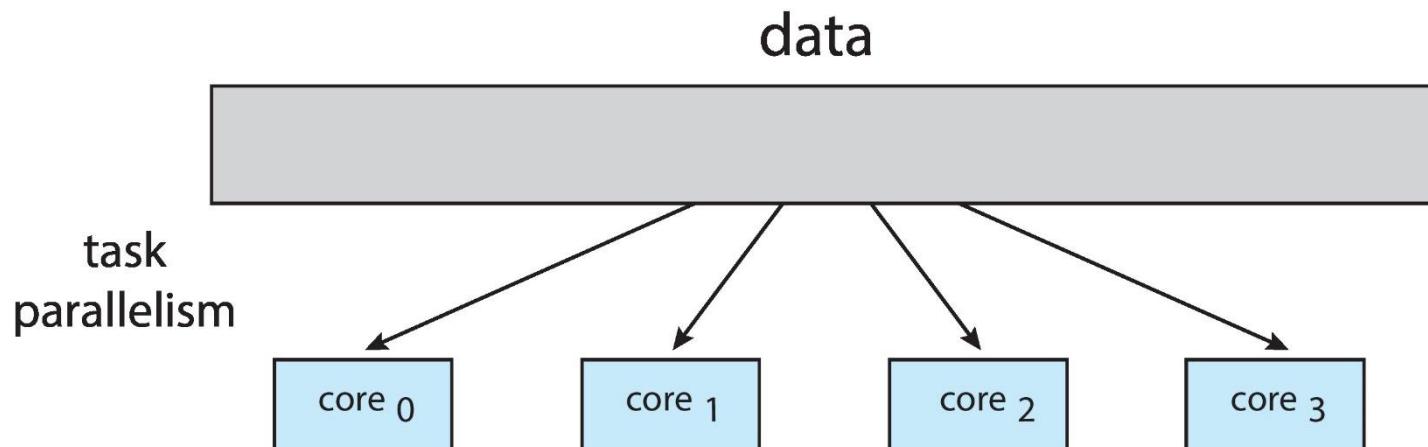
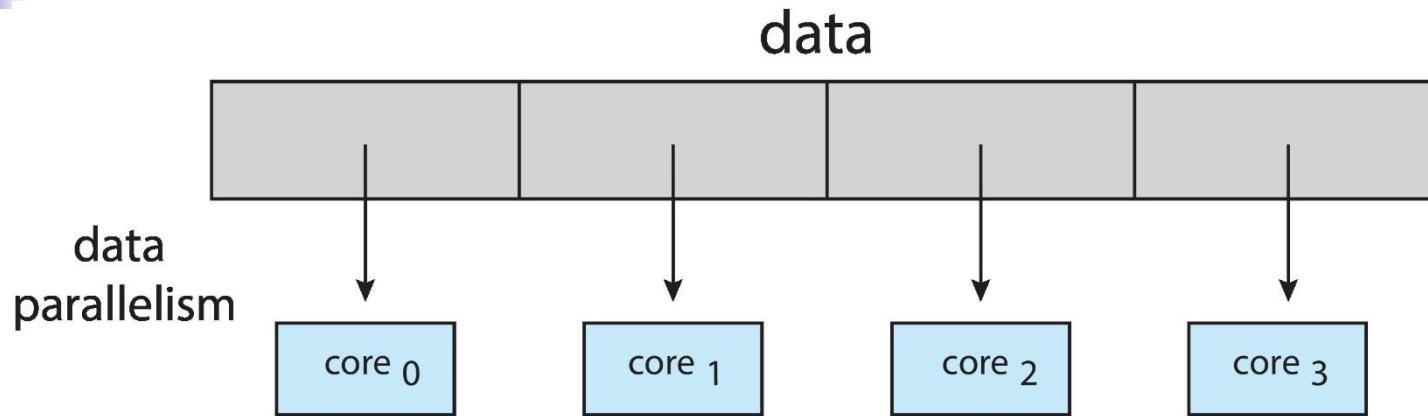
- **Task parallelism** – distributing tasks (threads) across multiple computing cores, each thread is performing a unique operation

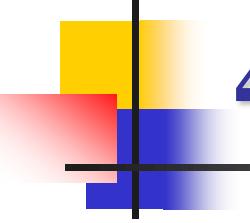
- e.g., thread A sums, thread B multiplies
- e.g., producer thread, consumer thread





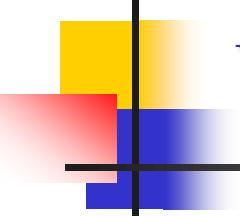
Data and task parallelism (Fig. 4.5)



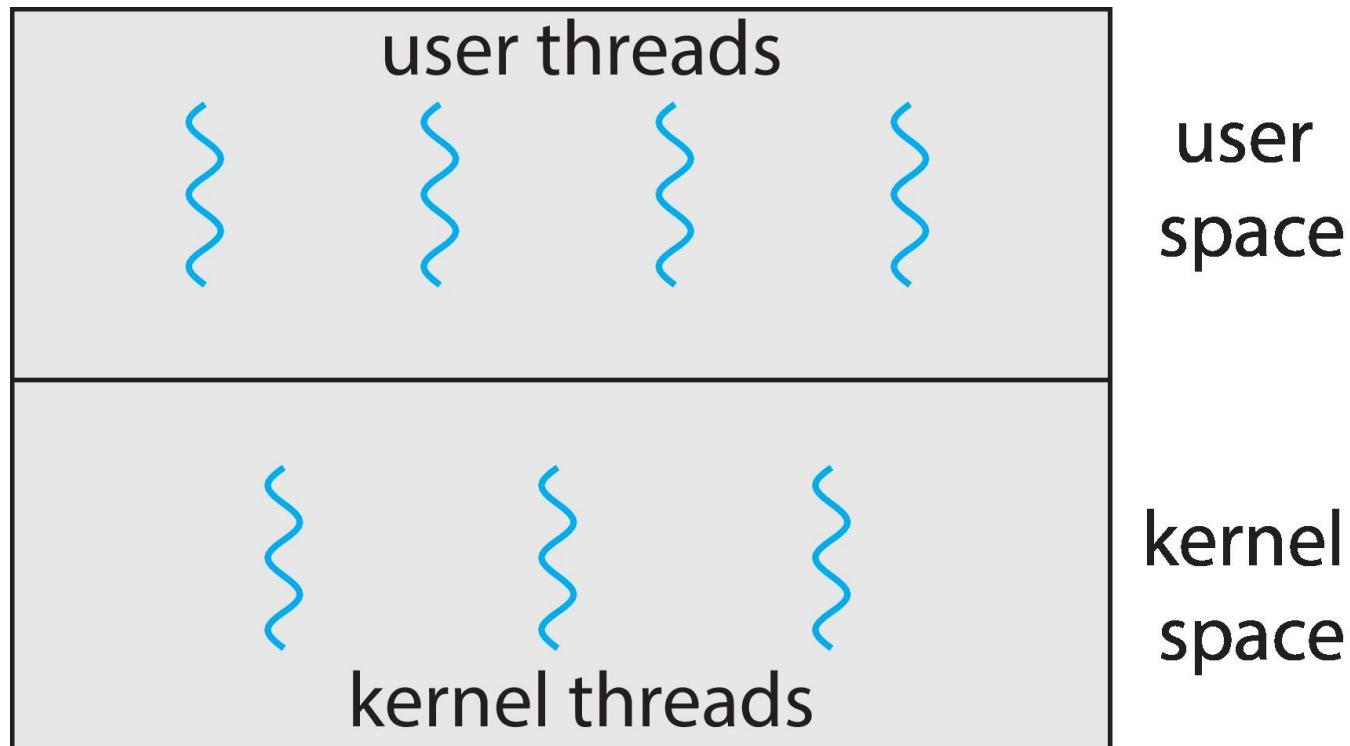


4.3 Multithreading Models

- Support for threads may be provided either at the user level, for user threads, or by the kernel, for kernel threads.
- **User threads** are supported above the kernel and are managed without kernel support, whereas **kernel threads** are supported and managed directly by the OS.
- Virtually all contemporary OSes support kernel threads:
 - Windows, Linux, macOS



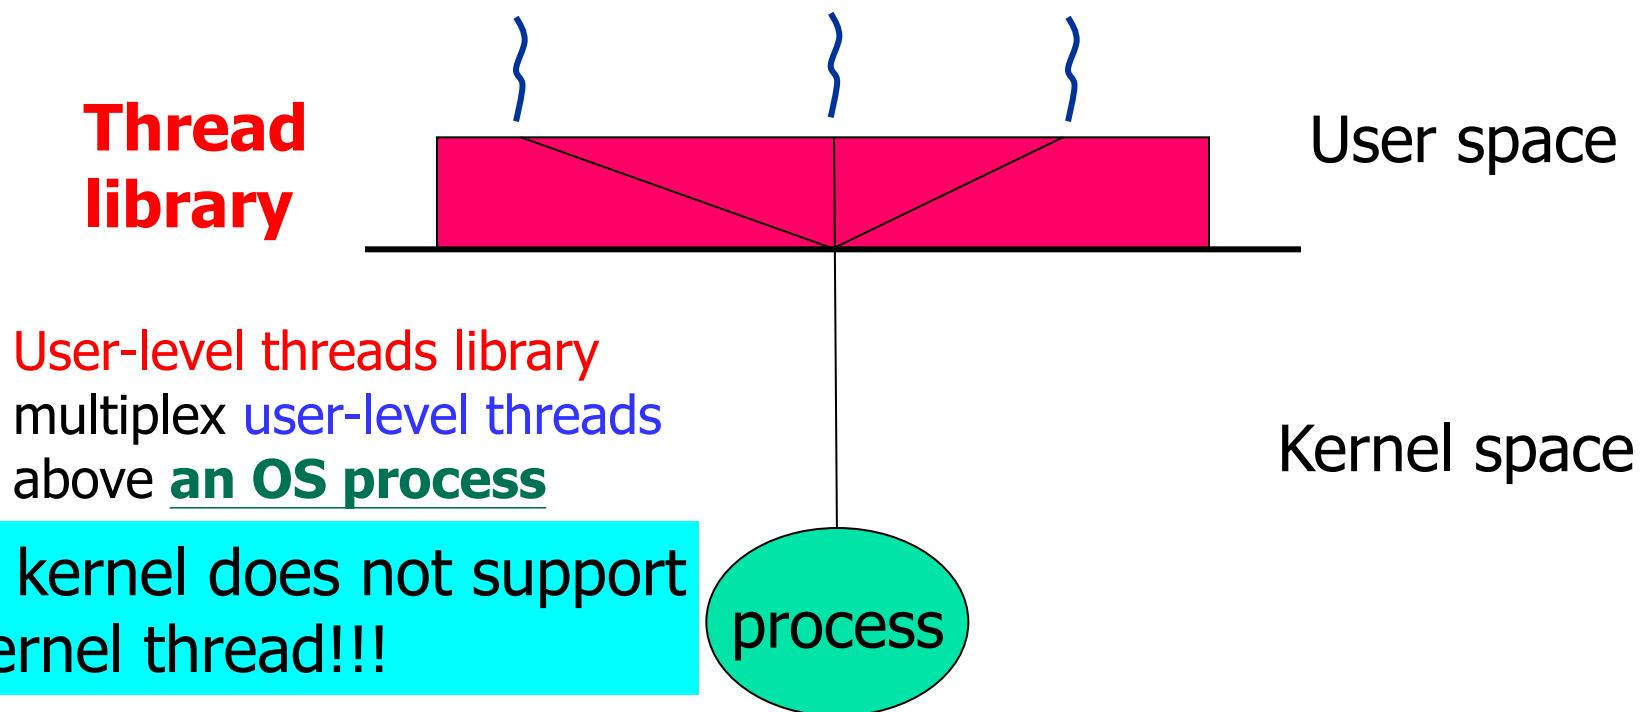
User and kernel threads (Fig. 4.6)



User-level Threads (1)

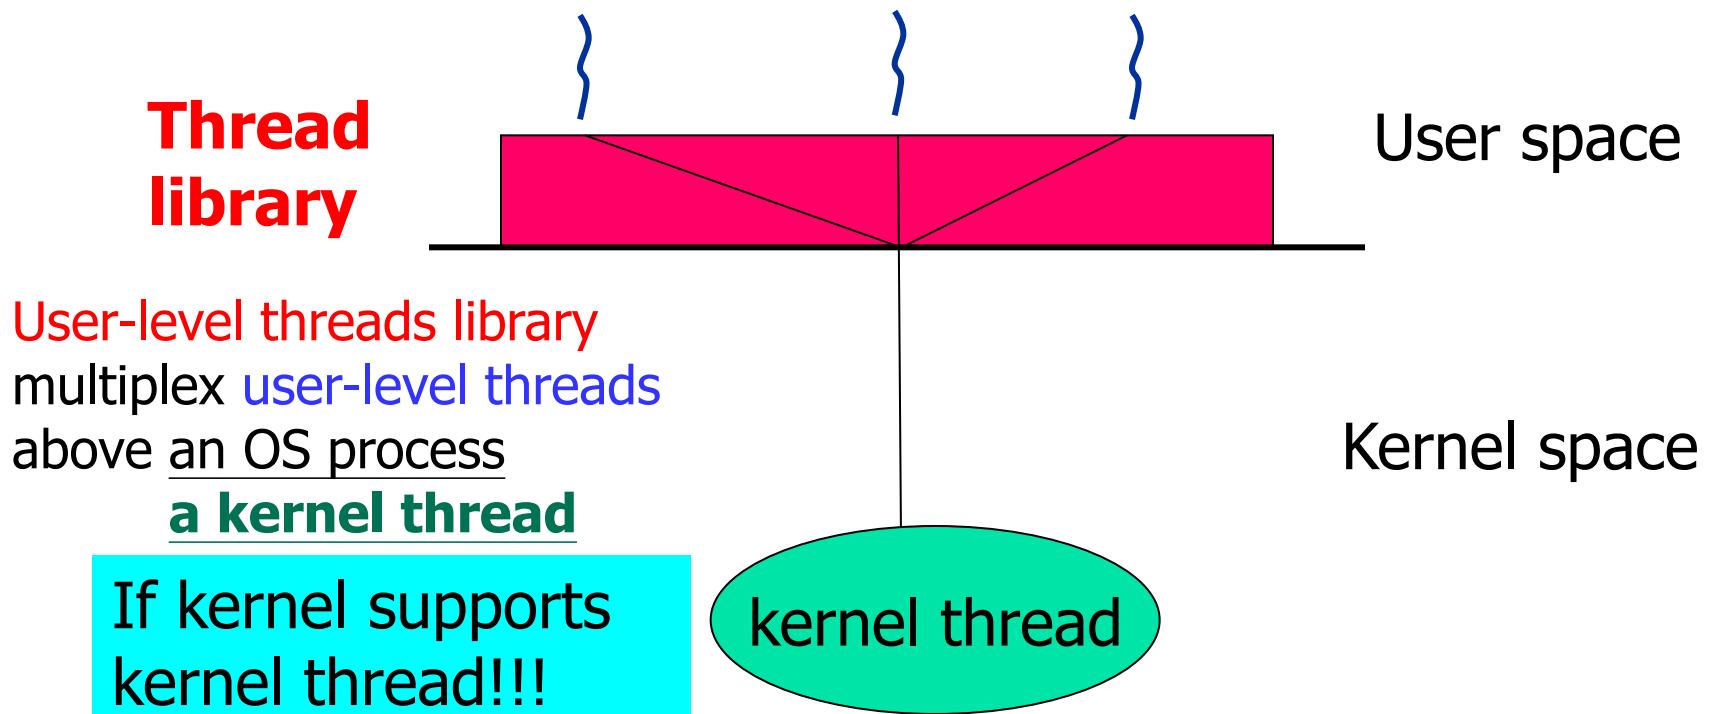
User-level threads are generally faster to create and manage than are **kernel threads**

- Because no intervention from the kernel is required
介入



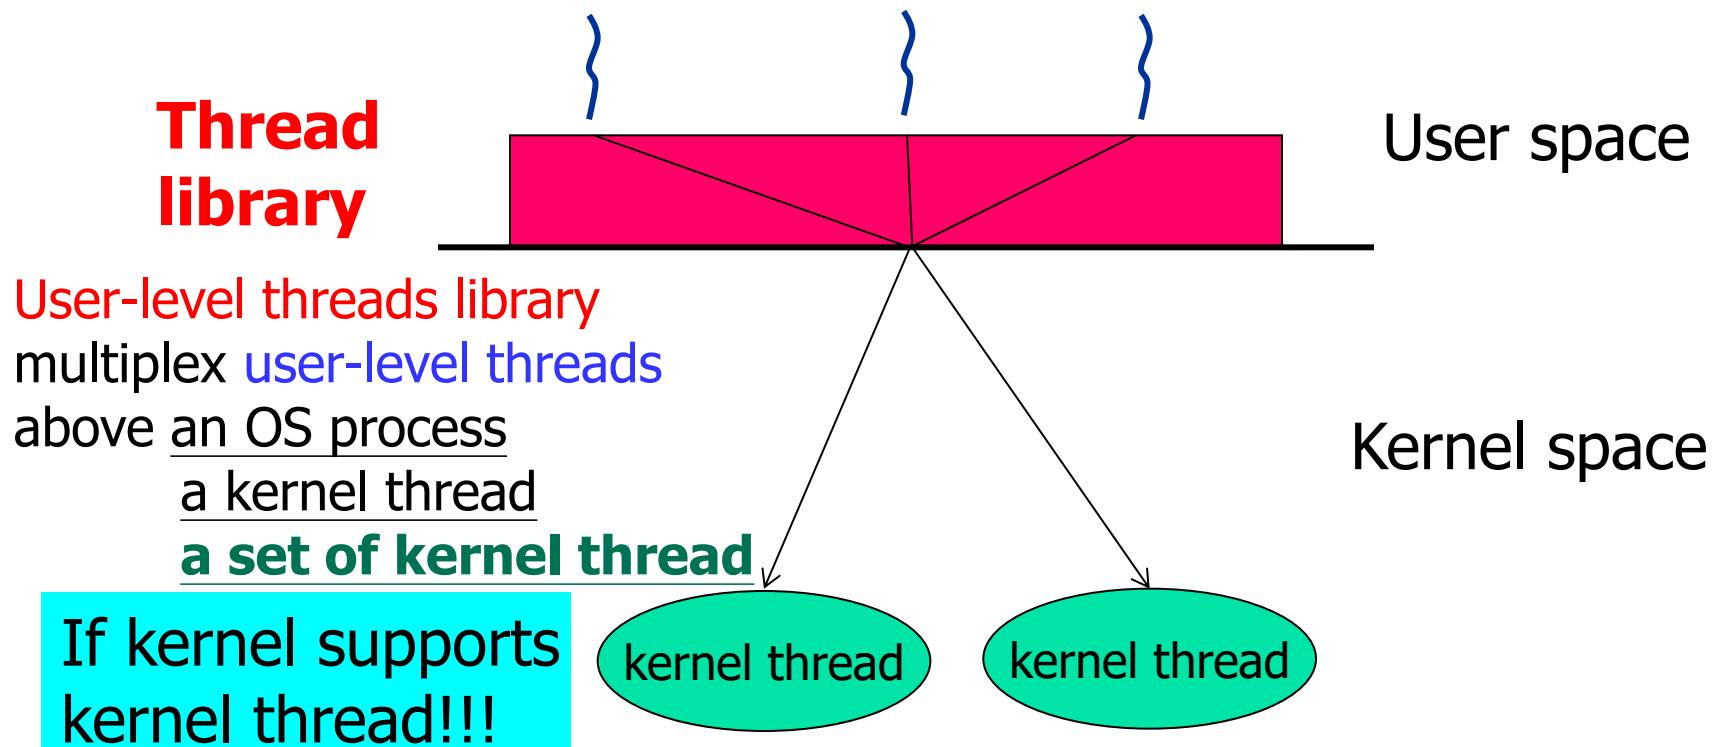
User-level Threads (2)

- **User-level threads** are generally faster to create and manage than are **kernel threads**
 - Because no intervention from the kernel is required

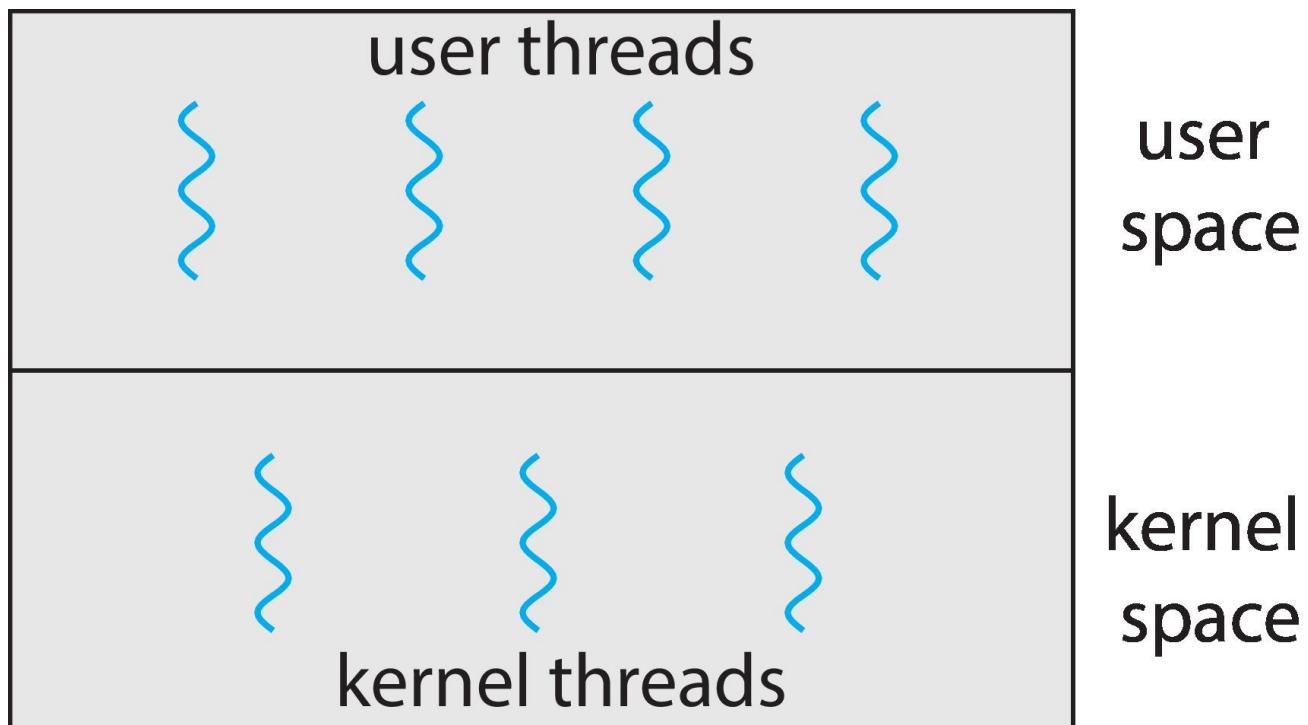


User-level Threads (3)

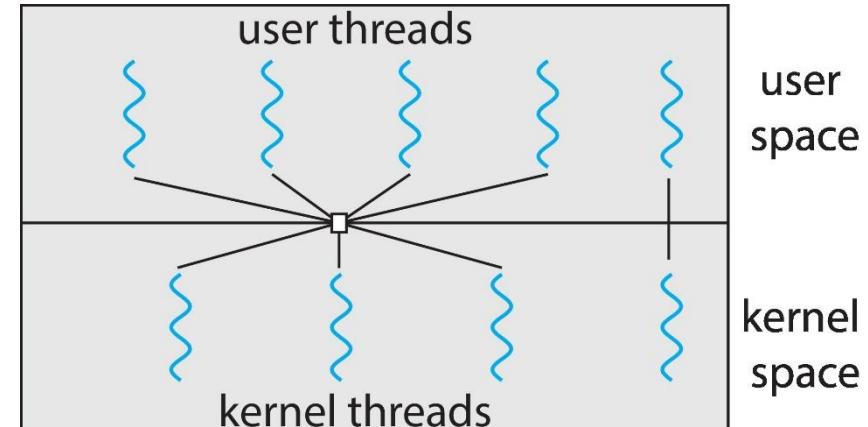
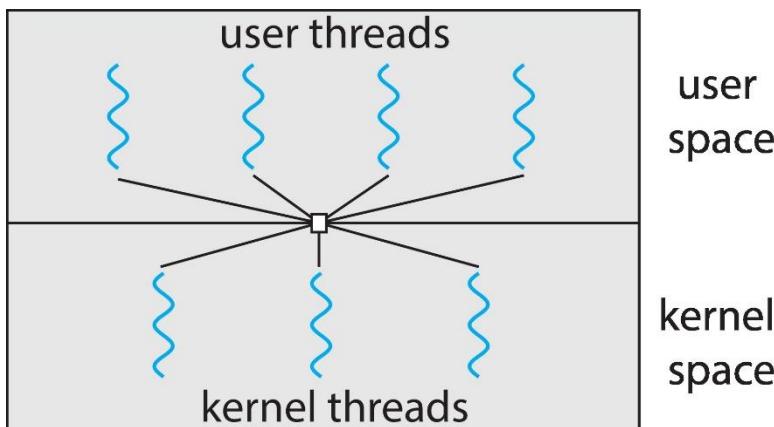
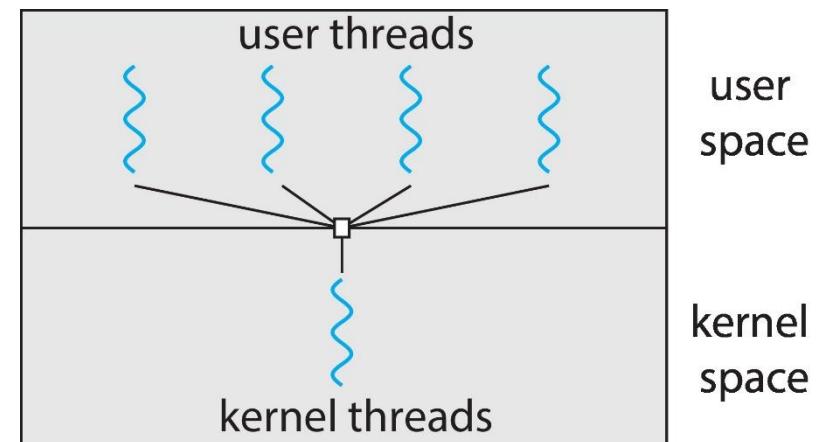
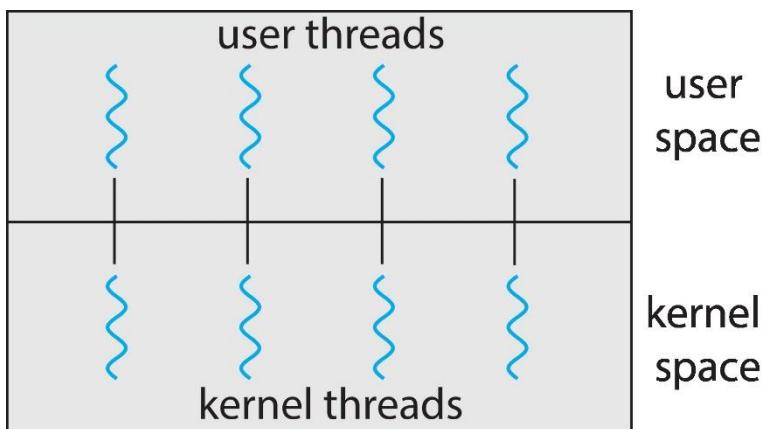
- **User-level threads** are generally faster to create and manage than are **kernel threads**
 - Because no intervention from the kernel is required

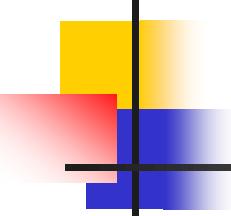


Relationship between user threads and kernel threads



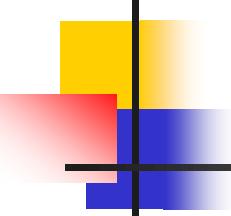
Relationship between user threads and kernel threads (Cont.)





Multithreading Models (Cont.)

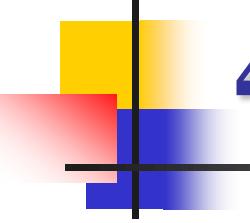
- Can threads run in parallel on multicore systems?
- If a thread makes a blocking system call, is the entire process still runnable?



Multithreading Models (Cont.)

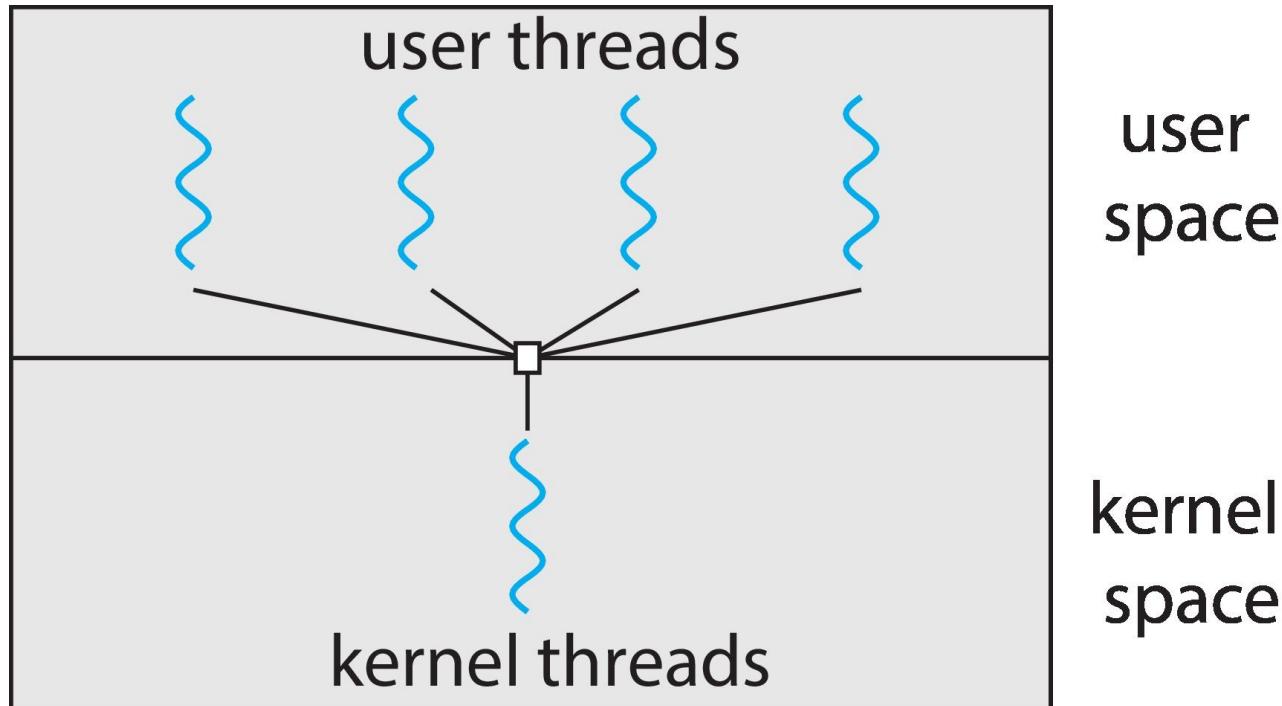
- Three common ways of establishing relationship between **user threads** and **kernel threads**:
 - 4.3.1 Many-to-One Model 
 - 4.3.2 One-to-One Model 
 - 4.3.3 Many-to-Many Model 

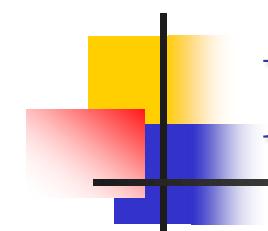




4.3.1 Many-to-One Model

- Maps many **user-level threads** to one **kernel thread**.





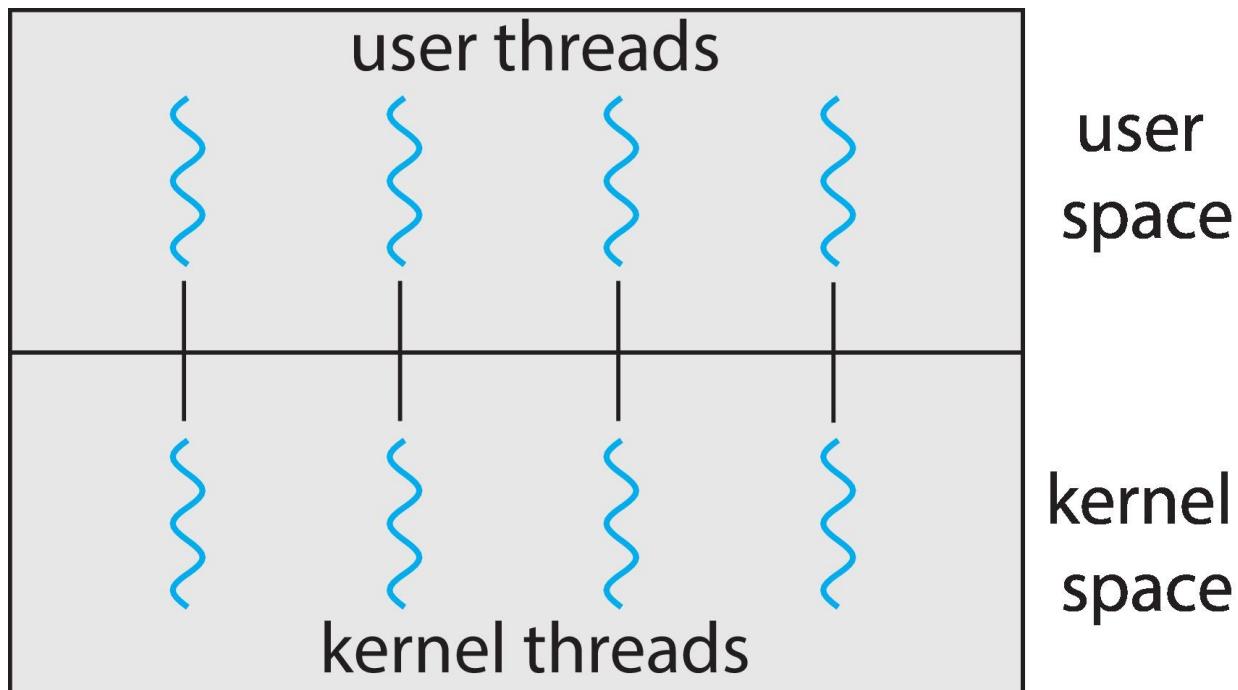
Many-to-One Model (Cont.)

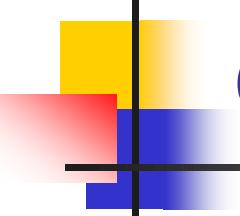
- Thread management is done by the thread library in user space, so it is efficient
- The entire process will block if a thread makes a blocking system call.
 - because only one thread can access the kernel at a time
 - multiple threads are unable to run in parallel on multicore systems
- Cannot take advantage of multiple processing cores
 - very few systems currently use this model
- Example:
 - Green threads – a thread library for Solaris systems



4.3.2 One-to-One Model

- Maps each **user thread** to a **kernel thread**.





One-to-One Model (Cont.)

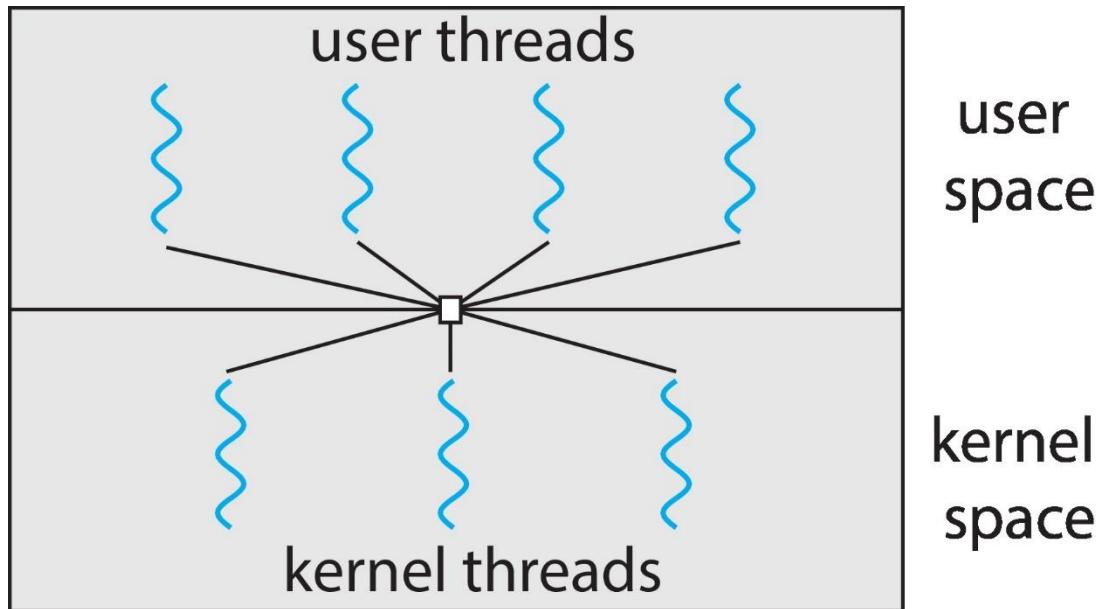
- Provides more concurrency than many-to-one model by allowing another thread to run when a thread makes a blocking system call.
- Allows multiple threads to run in parallel on multiprocessors.
- Creating a user thread requires creating the corresponding kernel thread, and a large number of kernel threads may burden the performance of a system.
- Examples
 - Linux, the family of Windows OSes

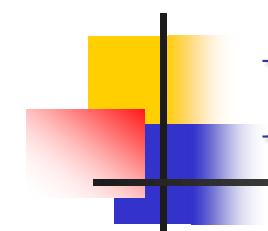


4.3.3 Many-to-Many Model

- Multiplexes many **user-level threads** to a smaller or equal number of **kernel threads**.

$$|\text{user threads}| \geq |\text{kernel threads}|$$





Discussion of Concurrency

■ **many-to-one model**

- Allows the developer to create as many user threads as she wishes, it does not result in parallelism, because the kernel can schedule only one kernel thread at a time.

■ **one-to-one model**

- Allows greater concurrency, but the developer has to be careful not to create too many threads within an application.
- When a thread performs a blocking system call, the kernel can schedule another thread for execution.

■ **many-to-many model**

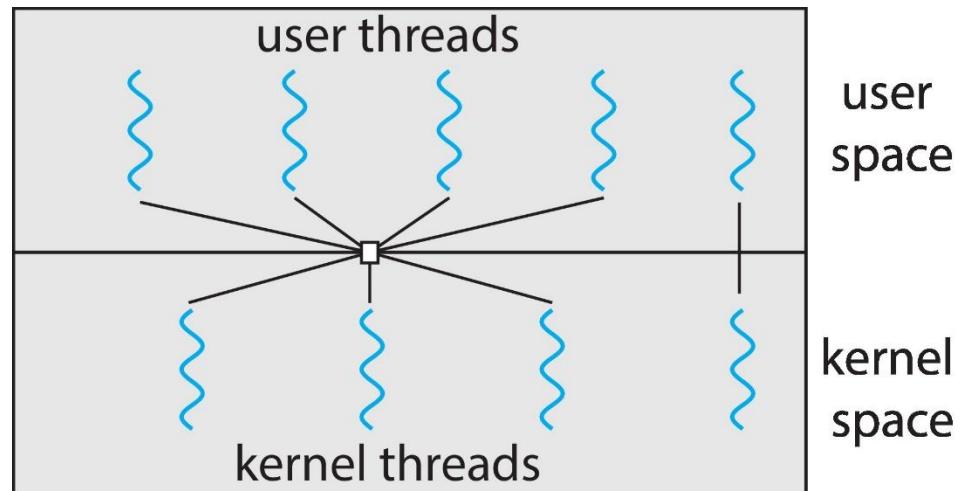
- Developers can create as many user threads as necessary, and the corresponding kernel threads can run in parallel on a multiprocessor.
- When a thread performs a blocking system call, the kernel can schedule another thread for execution.

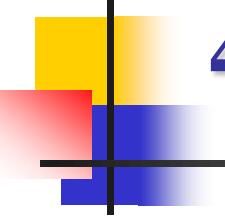
Two-level Model

變種

- **One variation on the many-to-many model**

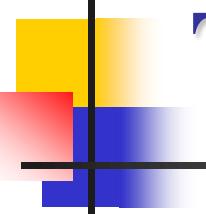
- still multiplexes many **user-level thread** to a smaller or equal number of **kernel threads**
- also allows a **user-level thread** to be bound to a **kernel thread**
- Appears to be the most flexible one, but is difficult to implement
- **With an increasing number of processing cores, limiting the number of kernel threads has become less important.**
 - Most OSes now use one-to-one model





4.4 Thread Libraries

- A **thread library** provides the programmer an API for creating and managing threads.
- Two primary ways of implementing a thread library:
 - (1) Provide a library entirely in user space with no kernel support
 - All code and data structures for the library exist in user space.
 - Invoking a function in the library results in a local function call in user space and not a system call
- (2) Implement a kernel-level library supported directly by OS
 - code and data structures for the library exist in kernel space.
 - Invoking a function in the API for the library results in a system call to the kernel



Thread Libraries (Cont.)

- Three main **thread libraries** are in use today:

- (1) **POSIX Pthreads**



- The threads extension of the POSIX standard, may be provided as either a user-level or kernel-level library.

- (2) **Windows**



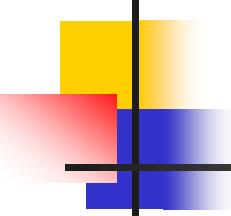
- A kernel-level library available on Windows systems

- (3) **Java**



- The **Java thread API** allows thread creation and management directly in Java programs.
 - The Java thread API is typically implemented using a thread library available on the host system.
 - Windows API on Windows systems
 - Pthreads API on UNIX, Linux, macOS





4.4.1 Pthreads

- **Pthreads refers to the POSIX standard (IEEE 1003.1c) defining an API for thread creation and synchronization.**
- ***Specification, not implementation***
- **API specifies behavior of the thread library, implementation is up to development of the library.**
- **Common in UNIX-type systems:**
 - Linux, macOS

Multithreaded C program using the Pthread API (Fig. 4.11)

```
/*
 * A pthread program illustrating how to
 * create a simple thread and some of the pthread API
 * This program implements the summation function where
 * the summation operation is run as a separate thread.
 * Usage on Solaris/Linux/Mac OS X:
 *
 * gcc thrd.c -lpthread
 * a.out <number>
```

compile & link with pthread lib.

```
#include <pthread.h>
#include <stdio.h>
```

```
int sum; /* this data is shared by the thread(s) */
```

```
void *runner(void *param); /* the thread */
```

$$sum = \sum_{i=1}^N i$$

```
main(int argc, char *argv[])
{
    pthread_t tid;          /* the thread identifier */
    pthread_attr_t attr;    /* set of attributes for the thread */
```

```
/* set the default attributes of the thread */
pthread_attr_init(&attr);
```

thread entry func.

```
/* create the thread */
```

```
pthread_create(&tid, &attr, runner, argv[1]);
```

```
/* now wait for the thread to exit */
pthread_join(tid, NULL);
```

argument

```
printf("sum = %d\n", sum);
```

```
}
```

```
/* The thread will execute in this function */
```

```
void *runner(void *param)
```

```
{
```

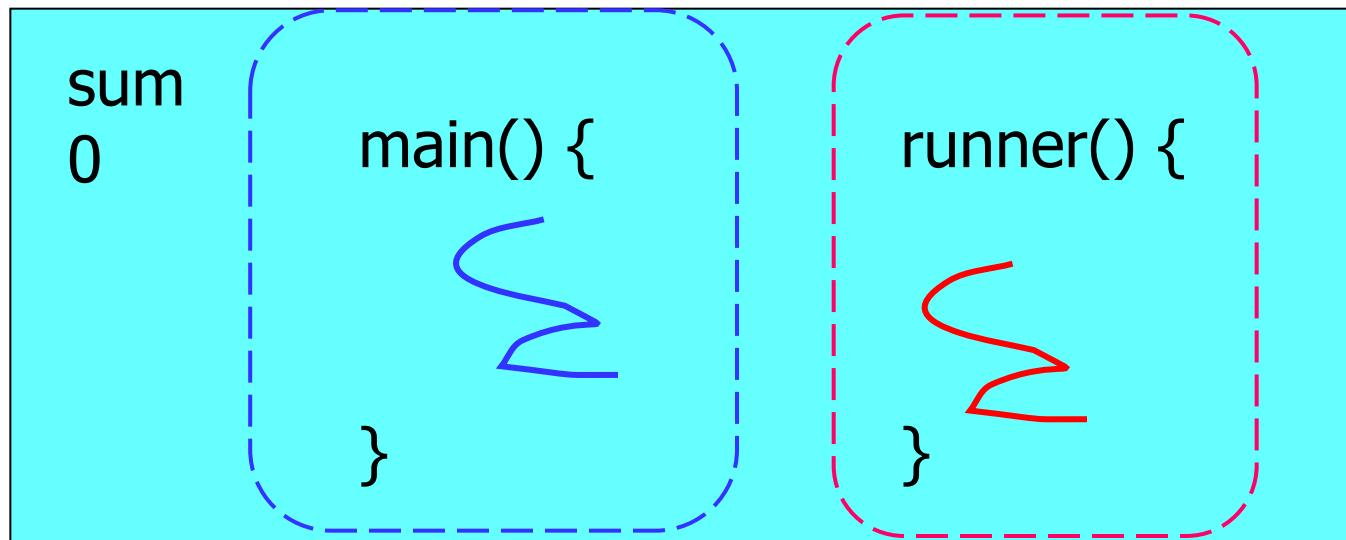
```
    int i, upper = atoi(param);
    sum = 0;
```

```
    for (i = 1; i <= upper; i++)
        sum += i;
```

```
    pthread_exit(0);
}
```

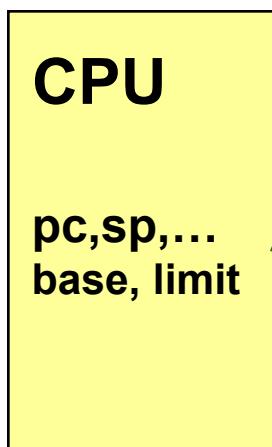
main thread

runner thread

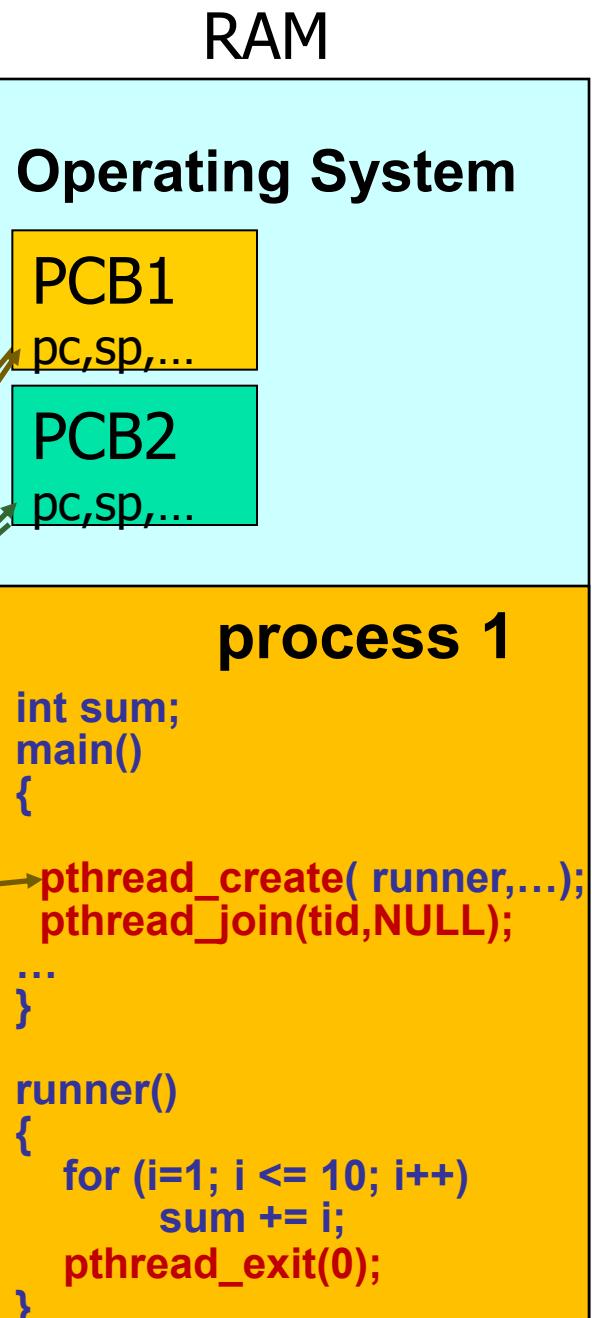


pthread_create()

CPU Scheduling



* if thread implementation with kernel support
e.g., Linux one-to-one mapping



Pthreads Code for Joining 10 Threads (Fig. 4.12)

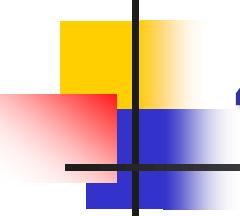
```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

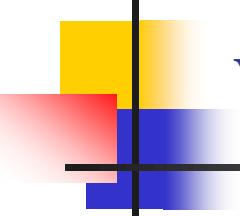
Figure 4.12 Pthread code for joining ten threads.





4.4.2 Windows Threads

- The technique for creating threads using the Windows thread library is similar to the Pthreads technique.
- Figure 4.13: Multithreaded C program using the Windows API.



Multithreaded C program using the Windows API (Fig. 4.13)

```
#include <stdio.h>
#include <windows.h>

DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD *)Param;

    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}
```

$$sum = \sum_{i=0}^N i$$

Multithreaded C program using Win32 API (Fig. 4.13) (Cont.)

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

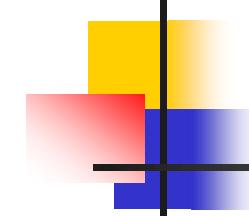
    Param = atoi(argv[1]);
    // create the thread
    ThreadHandle = CreateThread(NULL, 0, Summation,
                                &Param, 0, &ThreadId);

    // now wait for the thread to finish
    WaitForSingleObject(ThreadHandle, INFINITE);

    // close the thread handle
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```





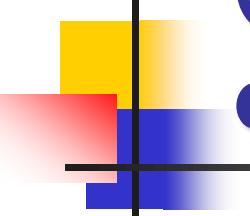
4.4.3 Java Threads

- Java provides support at the language level for the creation and management of threads.
- Java threads are managed by Java Virtual Machine.
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by:
 - (1) Extending Thread class 
 - (2) Defining a class that implements the Runnable interface 

```
public interface Runnable
{
    public abstract void run();
}
```

Standard practice is to implement Runnable interface





(1) Java program for the summation of a non-negative integer

```
public class Summation extends Thread
```

```
{
```

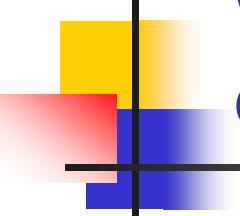
```
    public Summation(int n) {
        upper = n;
    }
```

```
    public void run() {
        int sum = 0;
        if (upper > 0) {
            for (int i = 1; i <= upper; i++)
                sum += i;
        }
    }
```

```
    System.out.println("The summation of " + upper + " is " + sum);
}
```

```
    private int upper;
}
```

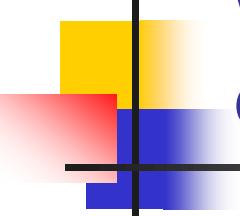
$$sum = \sum_{i=0}^N i$$



(1) Java program for the summation of a non-negative integer (Cont.)

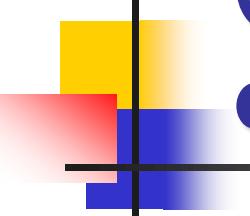
```
public class ThreadTester
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                throw new IllegalArgumentException(args[0] + "
must be non-negative.");
            else {
                Summation summationThread = new
                    Summation(Integer.parseInt(args[0]));
                summationThread.start();
            }
        } else
            System.err.println("Usage: Summation <integer value>");
    }
}
```





(2) Java program for the summation of a non-negative integer (Fig. 4.12)

```
/**  
 * This program creates a separate thread by implementing the Runnable interface.  
 * To compile, enter  
 * javac Driver.java  
 */  
class Sum  
{  
    private int sum;  
    public int getSum() {  
        return sum;  
    }  
    public void setSum(int sum) {  
        this.sum = sum;  
    }  
}
```



(2) Java program for the summation of a non-negative integer (Cont.)

```
class Summation implements Runnable
```

```
{
```

```
    private int upper;  
    private Sum sumValue;
```

```
    public Summation(int upper, Sum sumValue) {  
        if (upper < 0)  
            throw new IllegalArgumentException();  
        this.upper = upper;  
        this.sumValue = sumValue;
```

```
}
```

```
    public void run() {
```

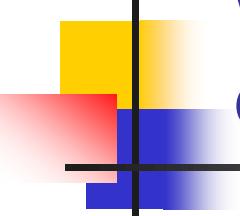
```
        int sum = 0;
```

```
        for (int i = 0; i <= upper; i++)  
            sum += i;
```

```
        sumValue.setSum(sum);
```

```
}
```

$$sum = \sum_{i=0}^N i$$

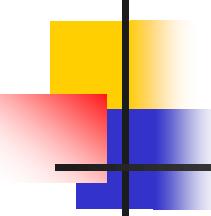


(2) Java program for the summation of a non-negative integer(Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length != 1) {
            System.err.println("Usage Driver <integer>");
            System.exit(0);
        }

        // create the object to be shared
        Sum sumObject = new Sum();
        int upper = Integer.parseInt(args[0]);

        Thread thrd = new Thread(new Summation(upper, sumObject));
        thrd.start();
        try {
            thrd.join();
        } catch (InterruptedException ie) { }
        System.out.println("The sum of " + upper + " is " + sumObject.get());
    }
}
```



Java Threads

Implementing Runnable interface:

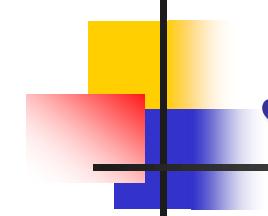
```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

Creating a thread:

```
Thread worker = new Thread(new Task());
worker.start();
```

Waiting on a thread:

```
try {
    worker.join();
}
catch (InterruptedException ie) { }
```



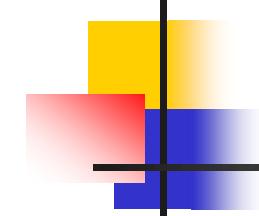
Java Executor Framework

- Rather than explicitly creating threads, Java also allows thread creation around the Executor interface:

```
public interface Executor
{
    void execute(Runnable command);
}
```

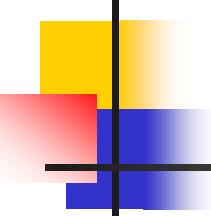
The Executor is used as follows:

```
Executor service = new Executor;
service.execute(new Task());
```



Java Executor Framework

```
import java.util.concurrent.*;  
  
class Summation implements Callable<Integer>  
{  
    private int upper;  
    public Summation(int upper) {  
        this.upper = upper;  
    }  
  
    /* The thread will execute in this method */  
    public Integer call() {  
        int sum = 0;  
        for (int i = 1; i <= upper; i++)  
            sum += i;  
  
        return new Integer(sum);  
    }  
}
```

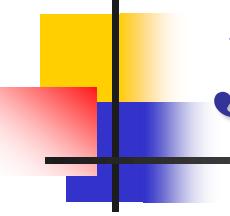


Java Executor Framework (cont.)

```
public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

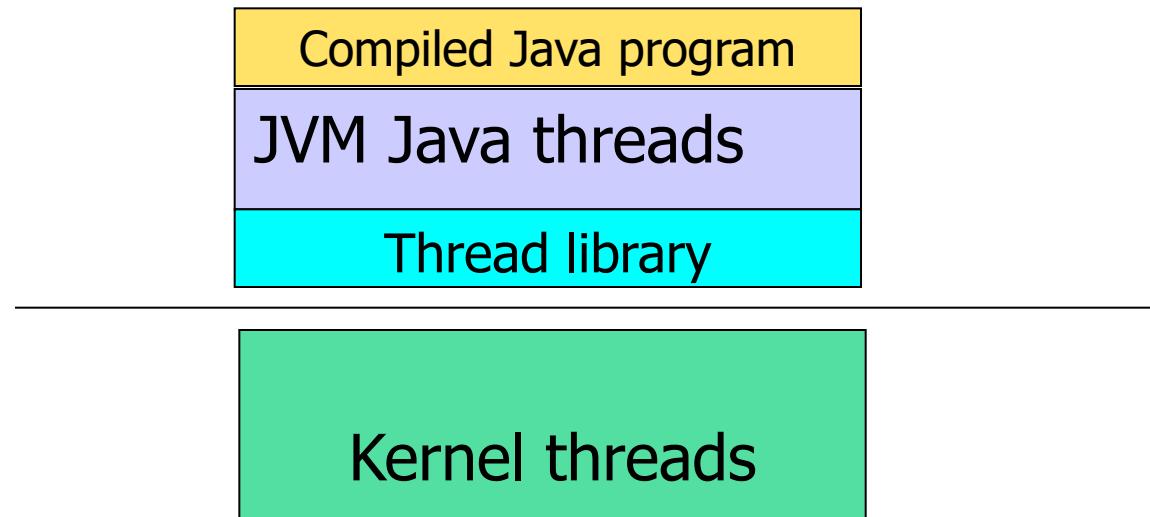
        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation(upper));

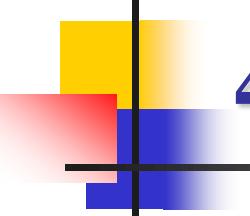
        try {
            System.out.println("sum = " + result.get());
        } catch (InterruptedException | ExecutionException ie) { }
    }
}
```



Java Threads (Cont.)

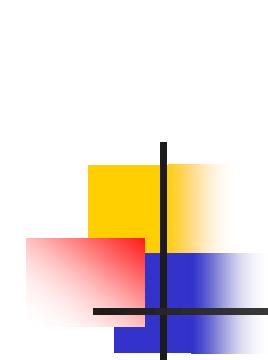
- The specification for the JVM does not indicate how **Java threads** are to be mapped to the underlying OS, instead leaving that decision to the particular implementation of JVM.





4.5 Implicit Threading

- **Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads**
- **Creation and management of threads done by compilers and run-time libraries rather than programmers**
- **Five methods explored**
 - Thread Pools
 - Fork-Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Threading Building Blocks



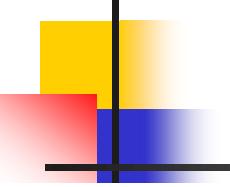
4.5.1 Thread pools

- **a multithreaded server**

- Whenever the server receives a request, it creates a separate thread to service the request.
- Potential problems:
 - thread creation time, unlimited threads could exhaust system resources

- **Thread pools**

- Create a number of threads at process startup and place them into a pool, where they sit and wait for work.
- When a server receives a request, it awakens a thread from this pool, passing it the request for service.
- Once the thread completes its service, it returns to the pool awaiting more work.



Thread pools (Cont.)

- **Benefits of thread pools**

- Faster to service a request with an existing thread than waiting to create a thread
- Limits the number of threads that exist at any one point.
- Separating task to be performed from mechanics of creating task allows different strategies for running task
 - i.e. tasks could be scheduled to run periodically

- **Examples:**

- Windows thread pool API

```
DWORD WINAPI PoolFunction(VOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```

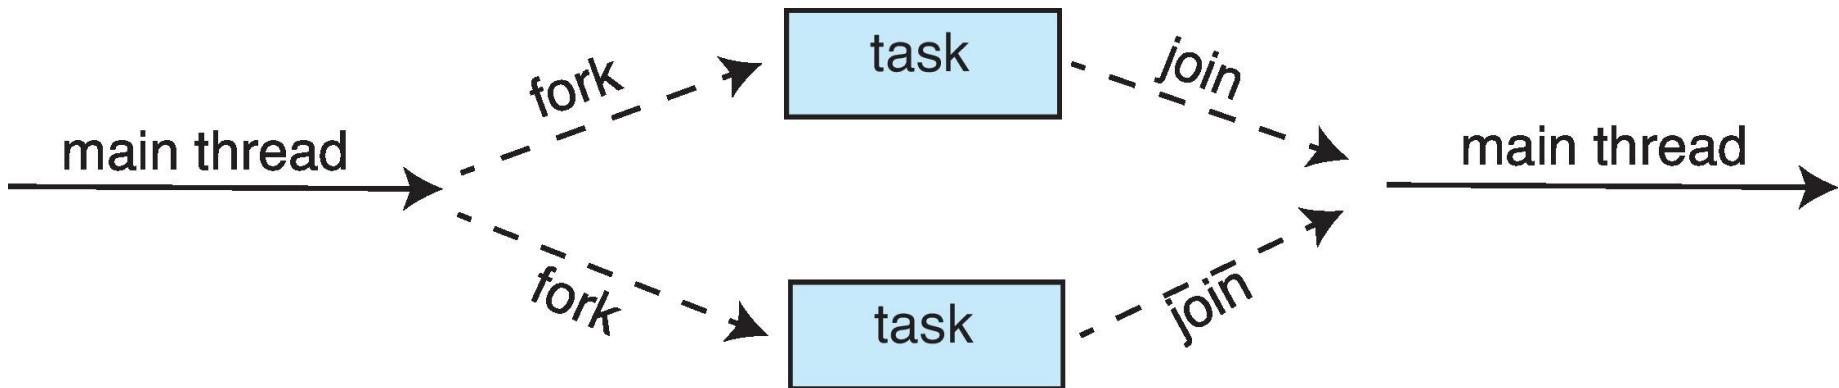
```
QueueUserWorkItem(&PoolFunction, NULL, 0);
```

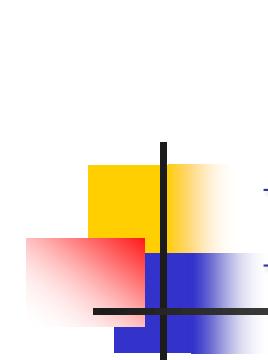
- Java provides a thread pool utility

4.5.2 Fork Join

Fork-Join Parallelism

- Multiple threads (tasks) are forked, and then joined.





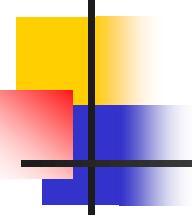
Fork-Join Parallelism

- General algorithm for fork-join strategy:

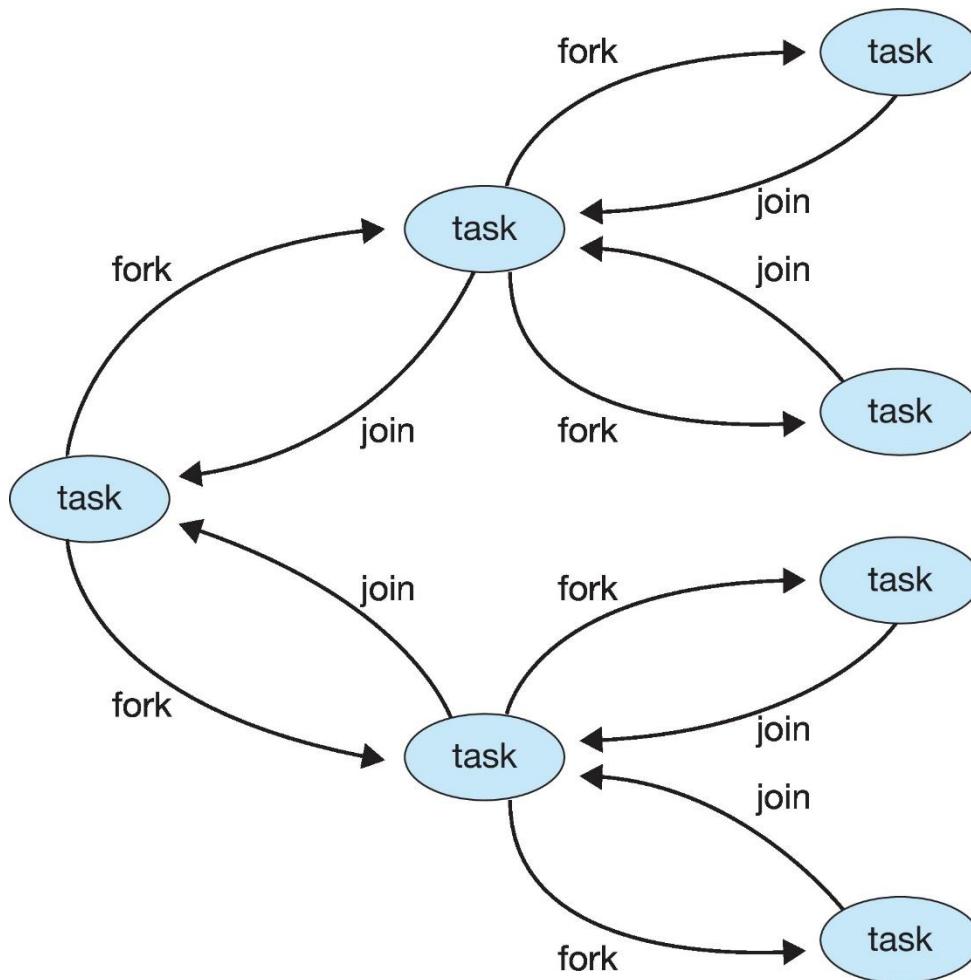
```
Task(problem)
    if problem is small enough
        solve the problem directly
    else
        subtask1 = fork(new Task(subset of problem)
        subtask2 = fork(new Task(subset of problem)

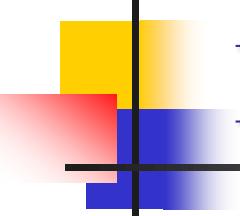
        result1 = join(subtask1)
        result2 = join(subtask2)

    return combined results
```



Fork-Join Parallelism





Fork-Join Parallelism in Java

```
ForkJoinPool pool = new ForkJoinPool();
// array contains the integers to be summed
int[] array = new int[SIZE];

SumTask task = new SumTask(0, SIZE - 1, array);
int sum = pool.invoke(task);
```



Fork-Join Parallelism in Java

```
import java.util.concurrent.*;

public class SumTask extends RecursiveTask<Integer>
{
    static final int THRESHOLD = 1000;

    private int begin;
    private int end;
    private int[] array;

    public SumTask(int begin, int end, int[] array) {
        this.begin = begin;
        this.end = end;
        this.array = array;
    }

    protected Integer compute() {
        if (end - begin < THRESHOLD) {
            int sum = 0;
            for (int i = begin; i <= end; i++)
                sum += array[i];

            return sum;
        }
        else {
            int mid = (begin + end) / 2;

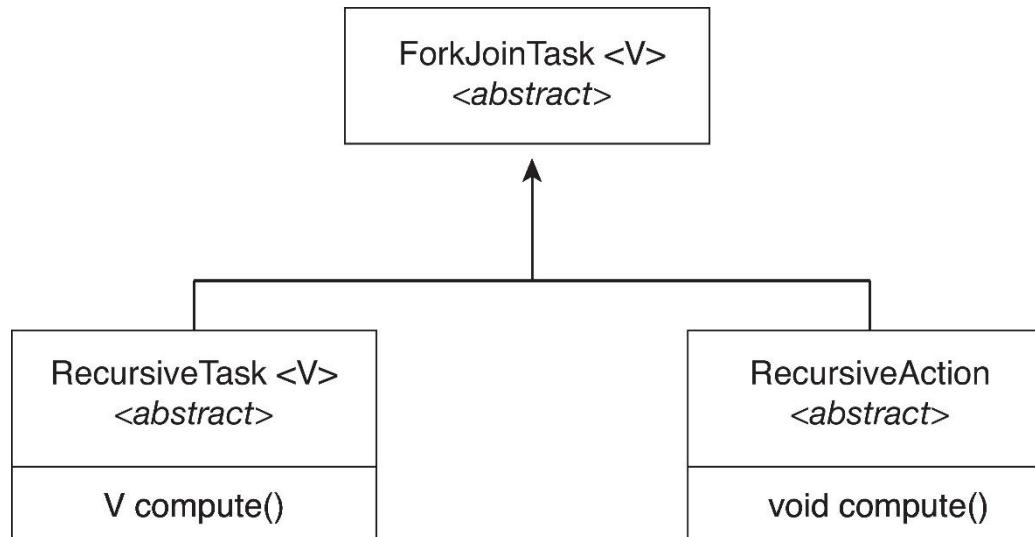
            SumTask leftTask = new SumTask(begin, mid, array);
            SumTask rightTask = new SumTask(mid + 1, end, array);

            leftTask.fork();
            rightTask.fork();

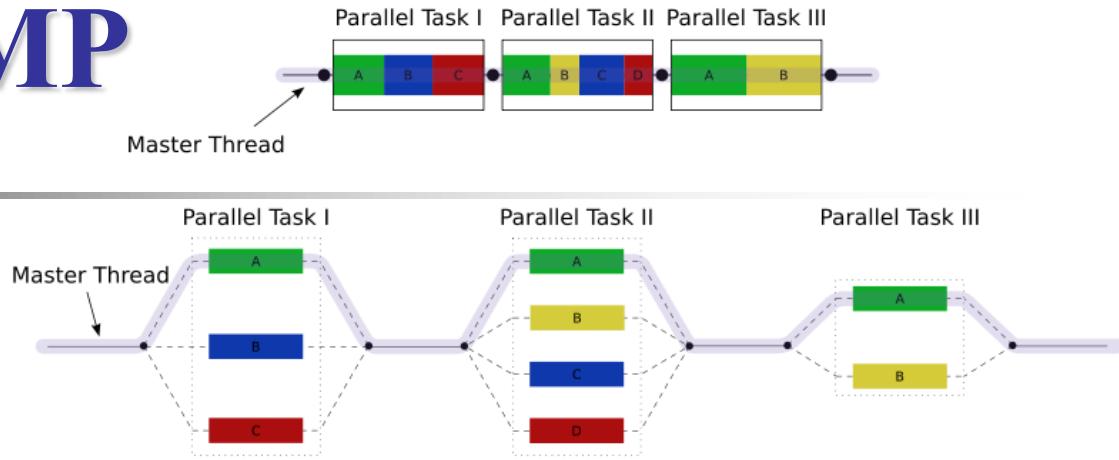
            return rightTask.join() + leftTask.join();
        }
    }
}
```

Fork-Join Parallelism in Java

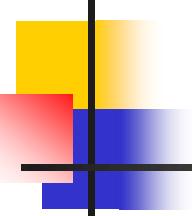
- The ForkJoinTask is an abstract base class
- RecursiveTask and RecursiveAction classes extend ForkJoinTask
- RecursiveTask returns a result (via the return value from the compute() method)
- RecursiveAction does not return a result



4.5.3 OpenMP



- A set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that may run in parallel
- OpenMP is available on several open-source and commercial compilers for Linux, Windows, and Mac OS X systems.



OpenMP (Cont.)

#pragma omp parallel

**Create as many threads
as there are cores**

#pragma omp parallel for

for(i=0;i<N;i++) {

c[i] = a[i] + b[i];

}

Run for loop in parallel

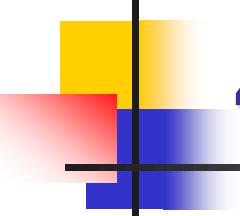
```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

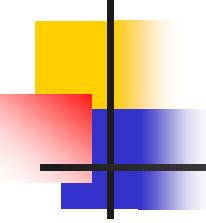
    /* sequential code */

    return 0;
}
```



4.5.4 Grand Central Dispatch

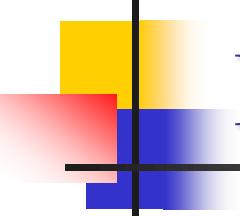
- A technology for Apple Mac OS X and iOS
- Extensions to C, C++ languages, API, and run-time library
- Allows application developers to identify parallel sections
- Manages most of the details of threading
- Block is in “`^{ }`” - `^{ printf("I am a block"); }`
- Blocks placed in dispatch queue
 - Assigned to available thread in thread pool when removed from queue



Grand Central Dispatch (Cont.)

- **Two types of dispatch queues:**
 - serial – blocks removed in FIFO order, queue is per process, called **main queue**
 - Programmers can create additional serial queues within program
 - concurrent – removed in FIFO order but several may be removed at a time
 - Three system wide queues with priorities low, default, high

```
dispatch_queue_t queue = dispatch_get_global_queue  
    (DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);  
  
dispatch_async(queue, ^{ printf("I am a block."); });
```



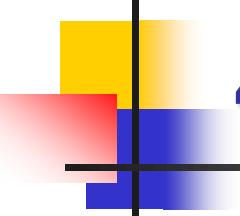
4.5.5 Intel Threading Building Blocks (TBB)

- **Template library for designing parallel C++ programs**
- **A serial version of a simple for loop**

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

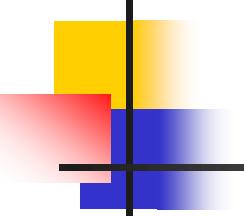
- **The same for loop written using TBB with parallel_for statement:**

```
parallel_for (size_t(0), n, [=](size_t i) {apply(v[i]);});
```



4.6 Threading Issues

- **4.6.1 fork() and exec() System Calls**
- **4.6.2 Signal Handling**
- **4.6.3 Thread Cancellation**
- **4.6.4 Thread-Local Storage**
- **4.6.5 Scheduler Activations**



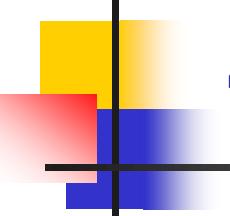
4.6.1 fork() & exec() System Calls

■ fork()

- What if one thread in a program calls fork():
 - Some UNIX systems have chosen to have two versions of fork():
 - The new process duplicates all threads
 - The new process is single-threaded (duplicates only the thread that invoked the fork())

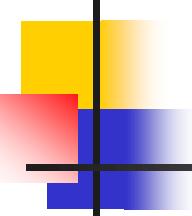
■ exec()

- What if a thread invokes the exec() system call?
 - The program specified in the parameter to exec() will replace the entire process – including all threads.



4.6.2 Signal Handling

- **signal**
 - Used in UNIX systems to notify a process that a particular event has occurred.
- **All signals follow the same pattern:**
 - A signal is generated by the occurrence of a particular event.
 - A generated signal is delivered to a process.
 - Once delivered, the signal must be handled.



Signal Handling (Cont.)

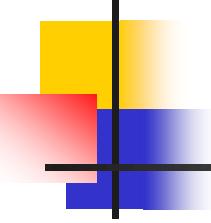
exception

- **Synchronous signal**

- A running program performs some actions, e.g., illegal memory access or division by zero, a signal is generated.
- Synchronous signals are delivered to the same process that performed the operation causing the signal

- **Asynchronous signal**

- A signal generated by an event external to a running process, that process receives the signal asynchronously.
- e.g., terminate a process with CTRL-C or timer expire
- Typically, an asynchronous signal is sent to another process.



Signal Handling (Cont.)

- **Every signal may be handled by one of two possible handlers:**
 - A default signal handler
 - that kernel runs when handling signal
 - A user-defined signal handler
 - can override default

```

#include <stdio.h>
#include <signal.h>

void sighup();
void sigint();
void sigquit();

int main(int argc, char *argv[]) {
    int pid;

    if ((pid = fork()) == 0) {
        printf("this is child process for receiving signal\n");
        signal(SIGHUP, sighup);
        signal(SIGINT, sigint);
        signal(SIGQUIT, sigquit);
        for(;;);
    } else {
        printf("this is parent process for sending signal\n");
        printf("\nPARENT: sending SIGHUP\n\n");
        kill(pid, SIGHUP);
        printf("\nPARENT: sending SIGINT\n\n");
        kill(pid, SIGINT);
        printf("\nPARENT: sending SIGQUIT\n\n");
        kill(pid, SIGQUIT);
        kill(pid, SIGHUP);
        kill(pid, SIGINT);
    }
}

void sighup() {
    printf("CHILD: I have received a SIGHUP\n");
}

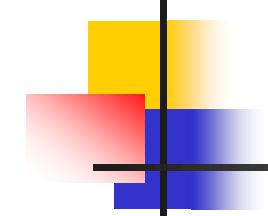
void sigint() {
    printf("CHILD: I have received a SIGINT\n");
}

void sigquit() {
    printf("My DADDY has Killed me!!!\n");
}

```

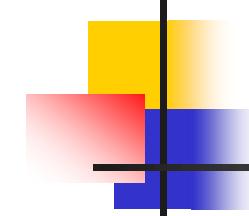
user-defined signal handlers

sending signals to process



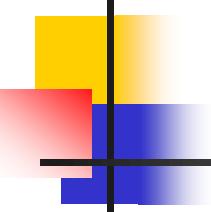
Signal Handling (Cont.)

- How to deal with signals in a multithreaded environment?
 - Deliver the signal to the thread to which the signal applies.
 - e.g., synchronous signals
 - Deliver the signal to every thread in the process.
 - e.g., CRTL-C
 - Deliver the signal to certain threads in the process.
 - e.g., Pthreads provides **pthread_kill(threadid, signal)**
 - Assign a specific thread to receive all signals for the process.
 - e.g., Solaris 2



4.6.3 Thread Cancellation

- **The task of terminating a thread before it has completed.**
- Thread to be canceled is **target thread**
- **Asynchronous cancellation**
 - One thread immediately terminates the target thread
- **Deferred cancellation**
 - The target thread periodically checks whether it should terminate
 - Allows a thread to check if it should be canceled at a point at which it can be canceled safely.
 - Example: Pthreads refers to such points as cancellation points.



Thread Cancellation

- Pthread code to create and cancel a thread:

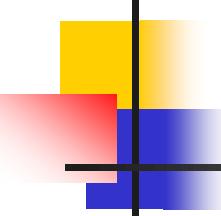
```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
.  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);  
  
/* wait for the thread to terminate */  
pthread_join(tid,NULL);
```

Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	-
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - i.e., `pthread_testcancel()`
 - Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals



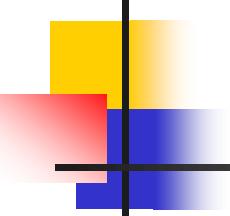
Thread Cancellation in Java

- Deferred cancellation uses the `interrupt()` method, which sets the interrupted status of a thread.

```
Thread worker;  
  
    . . .  
  
/* set the interruption status of the thread */  
worker.interrupt()
```

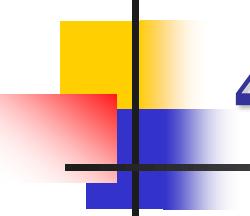
- A thread can then check to see if it has been interrupted:

```
while (!Thread.currentThread().isInterrupted()) {  
    . . .  
}
```



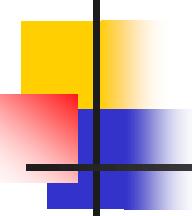
4.6.4 Thread-Local Storage

- Each thread might need its own copy of certain data in some circumstances – **thread-local storage (TLS)**.
 - Visible across function invocations
- Most thread libraries – including Windows and Pthreads – provide some form of support for thread-local storage.
- Java provides support as well.

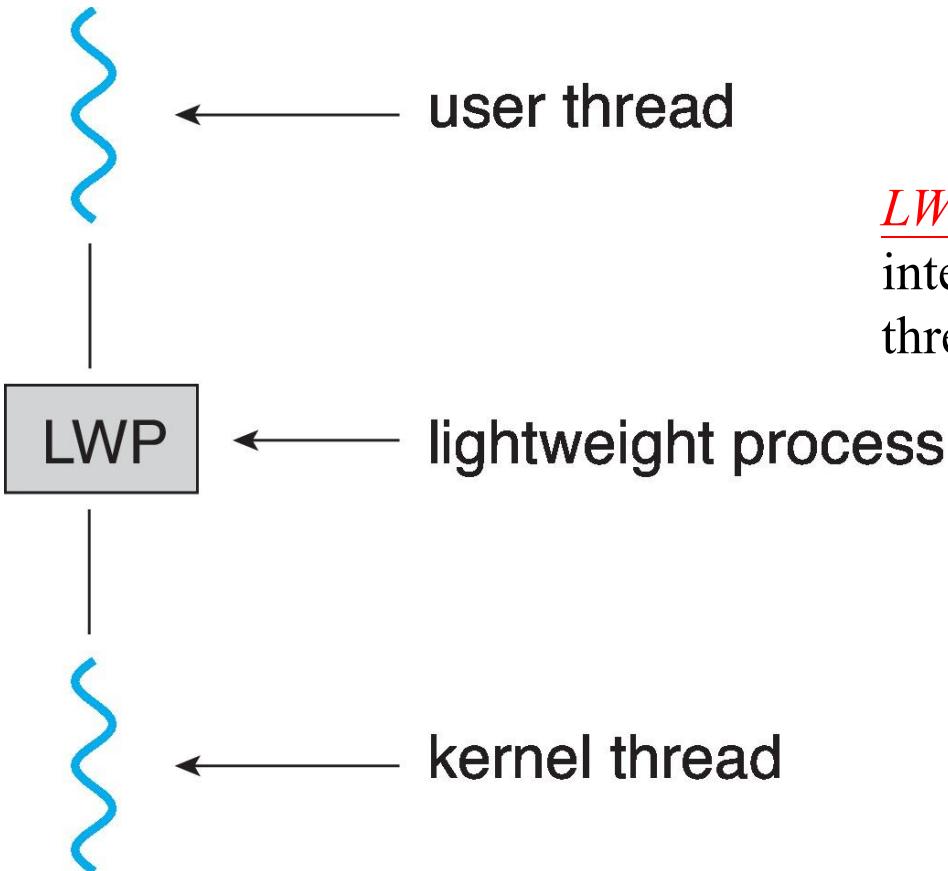


4.6.5 Scheduler Activations

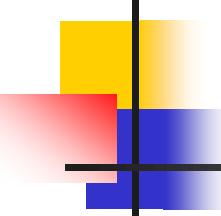
- **Concerns communication between the kernel and the thread library.**
- **Both many-to-many and two-level models require communication to maintain the appropriate number of kernel threads allocated to the application**
- **Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)****
 - Appears to be a virtual processor on which the application can schedule a user thread to run
 - Each LWP is attached to a kernel thread



Lightweight process (LWP)



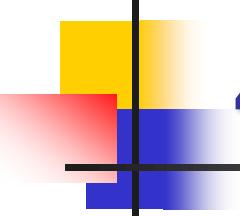
LWP (lightweight processes) - intermediate level between user-level threads and kernel-level threads.



Scheduler Activations (Cont.)

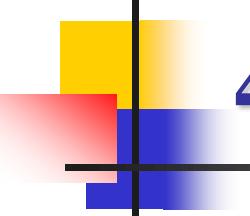
- **Scheduler activations provide upcalls - a communication mechanism from the kernel to the thread library**
 - Upcalls are handled by the thread library with an *upcall handler*
- **This communication allows an application to maintain the correct number kernel threads**





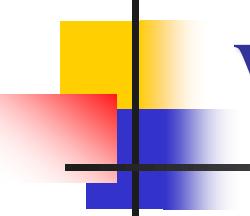
4.7 Operating-System Examples

- **Explore how threads are implemented in Windows and Linux systems.**
 - 4.7.1 Windows Threads
 - 4.7.2 Linux Threads



4.7.1 Windows Threads

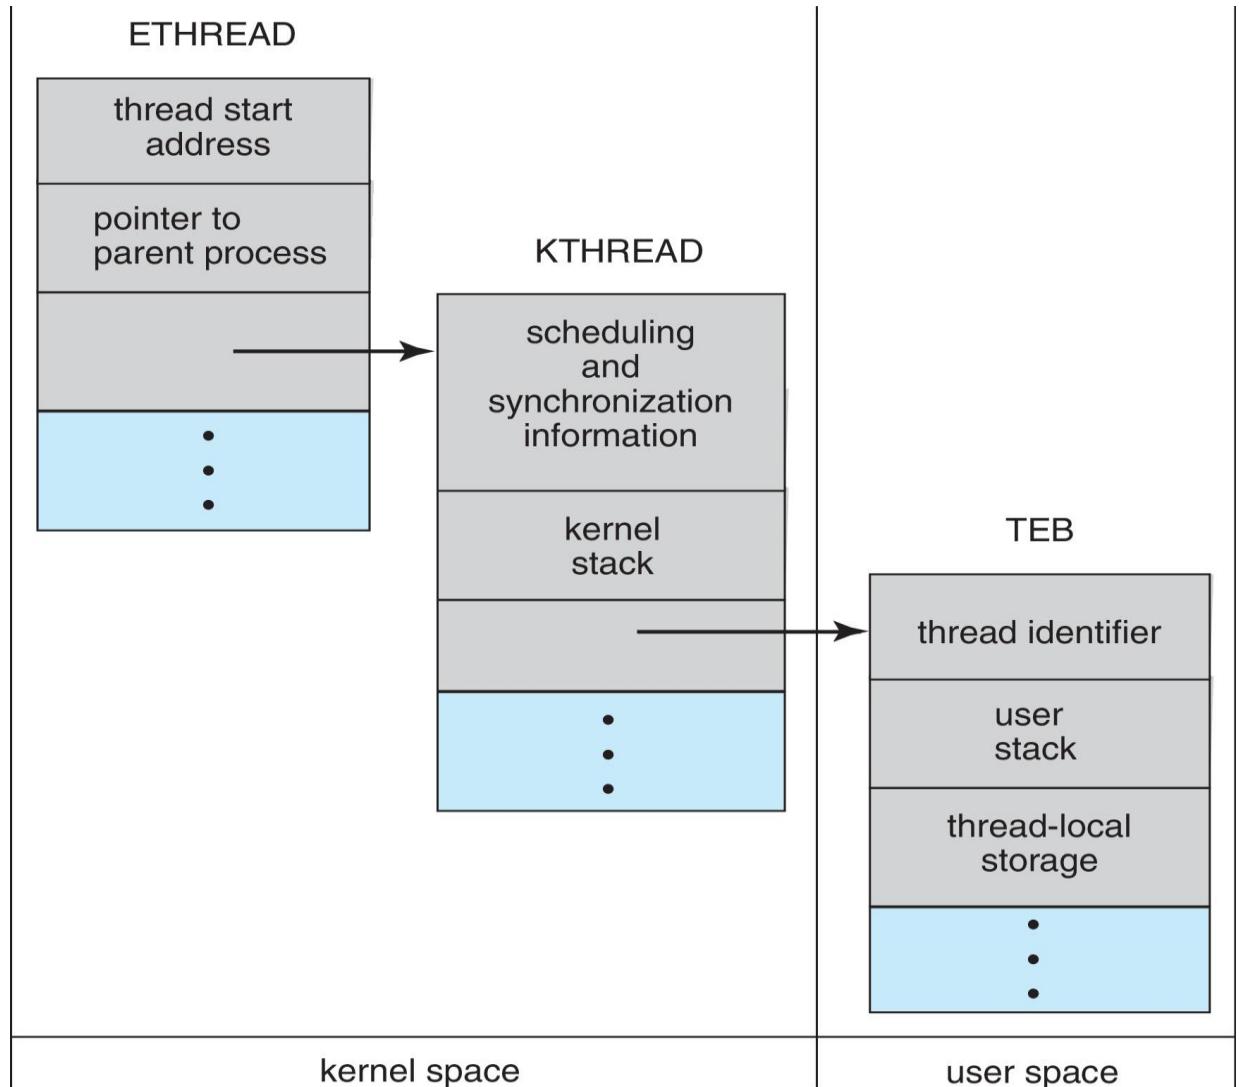
- Windows implements the Windows API.
- A Windows application runs as a separate process, and each process may contain one or more threads.
- Implements the one-to-one mapping, where each user-level thread maps to an associated kernel thread.

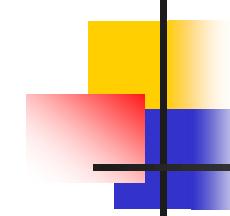


Windows Threads (Cont.)

- **Each thread contains**
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- **The register set, stacks, and private storage area are known as the context of the threads**
- **The primary data structures of a thread include:**
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

Data structures of a Windows thread



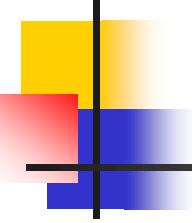


4.7.2 Linux Threads

- **Linux does not distinguish between processes and threads.**
 - Linux generally uses the term **tasks** rather than *process* or *threads* when referring to a flow of control within a program.
- **Thread creation is done through `clone()` system call.**
- **`clone()` allows a child task to share the address space of the parent task (process).**
 - the passed flags determine how much sharing between parent & child

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.





Homework

- **10th edition**

- Ex. 4.4, 4.5, 4.6, 4.9, 4.11, 4.17, 4.19

- Ex. 4.19 手寫 (40%) 或程式實作 (50%)

- if (手寫)
 - 最高得 40分
 - else
 - 最高得 50分
 - 請印出執行結果畫面 (使用自己的帳號和虛擬機).
 - 請交原始程式

EX 4.19

```
#include <pthread.h>
#include <stdio.h>

int value = 0;
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pid_t pid;
    pthread_t tid;
    pthread_attr_t attr;

    pid = fork();

    if (pid == 0) { /* child process */
        pthread_attr_init(&attr);
        pthread_create(&tid,&attr,runner,NULL);
        pthread_join(tid,NULL);
        printf("CHILD: value = %d",value); /* LINE C */
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE P */
    }
}

void *runner(void *param) {
    value = 5;
    pthread_exit(0);
}
```